

Oracle SQL tuning  
- An Overview

## 3 P's



### ➤ **Purpose :**

- To learn an overview of Performance Tuning

### ➤ **Product :**

- » Learn to understand the basics of Tuning
  - Describe what attributes of a SQL statement can make it perform poorly
  - List the Oracle tools that can be used to tune SQL

### ➤ **Process:**

- Instructor led training with practical experience



- » Introduction to SQL tuning
- » Describe why the SQL statements are performing poorly
- » Introduction to Oracle Optimizer
- » Discuss the need for Optimizer
- » Explain the various phases of Optimization
- » Gather Execution Plans
- » Interpret Execution Plans
- » Interpret the output of TKPROF
- » Gather Optimizer statistics
- » Use Hints appropriately



- » Reasons for Inefficient SQL Performance
- » SQL Tuning tasks
- » Proactive tuning Methodology



## Reasons for Inefficient SQL Performance

- Stale or missing optimizer statistics
- Missing access structures
- Suboptimal execution plan selection
- Poorly constructed SQL



```
SELECT COUNT(*) FROM products p WHERE prod_list_price < 1.15  
(SELECT avg(unit_cost) FROM costs c  
WHERE c.prod_id = p.prod_id)
```

```
SELECT * FROM job_history jh, employees e  
WHERE substr(to_char(e.employee_id),2) =  
substr(to_char(jh.employee_id),2)
```

```
SELECT * FROM orders WHERE order_id_char = 1205
```

```
SELECT * FROM employees  
WHERE to_char(salary) = :sal
```

```
SELECT * FROM parts_old  
UNION  
SELECT * FROM parts_new
```



## SQL Tuning Tasks: Overview

- Identifying high-load SQL
- Gathering statistics
- Generating system statistics
- Rebuilding existing indexes
- Maintaining execution plans
- Creating new index strategies



# Scalability with Application Design, Implementation, and Configuration

- Applications have a significant impact on scalability.
  - Poor schema design can cause expensive SQL that does not scale.
  - Poor transaction design can cause locking and serialization problems.
  - Poor connection management can cause unsatisfactory response times.





## Common Mistakes on Customer Systems

1. Bad connection management
2. Bad use of cursors and the shared pool
3. Excess of resources consuming SQL statements
4. Use of nonstandard initialization parameters
5. Poor database disk configuration
6. Redo log setup problems
7. Excessive serialization
8. Inappropriate full table scans
9. Large number of recursive SQL statements related to space management or parsing activity
10. Deployment and migration errors



# Proactive Tuning Methodology

- Simple design
- Data modeling
- Tables and indexes
- Using views
- Writing efficient SQL
- Cursor sharing
- Using bind variables



## Simplicity in Application Design

- Simple tables
- Well-written SQL
- Indexing only as required
- Retrieving only required information



# Table Design

- Compromise between flexibility and performance:
  - Principally normalize
  - Selectively denormalize
- Use Oracle performance and management features:
  - Default values
  - Constraints
  - Materialized views
  - Clusters
  - Partitioning
- Focus on business-critical tables

# Index Design



- Create indexes on the following:
  - Primary key (can be automatically created)
  - Unique key (can be automatically created)
  - Foreign keys (good candidates)
- Index data that is frequently queried (select list).
- Use SQL as a guide to index design.



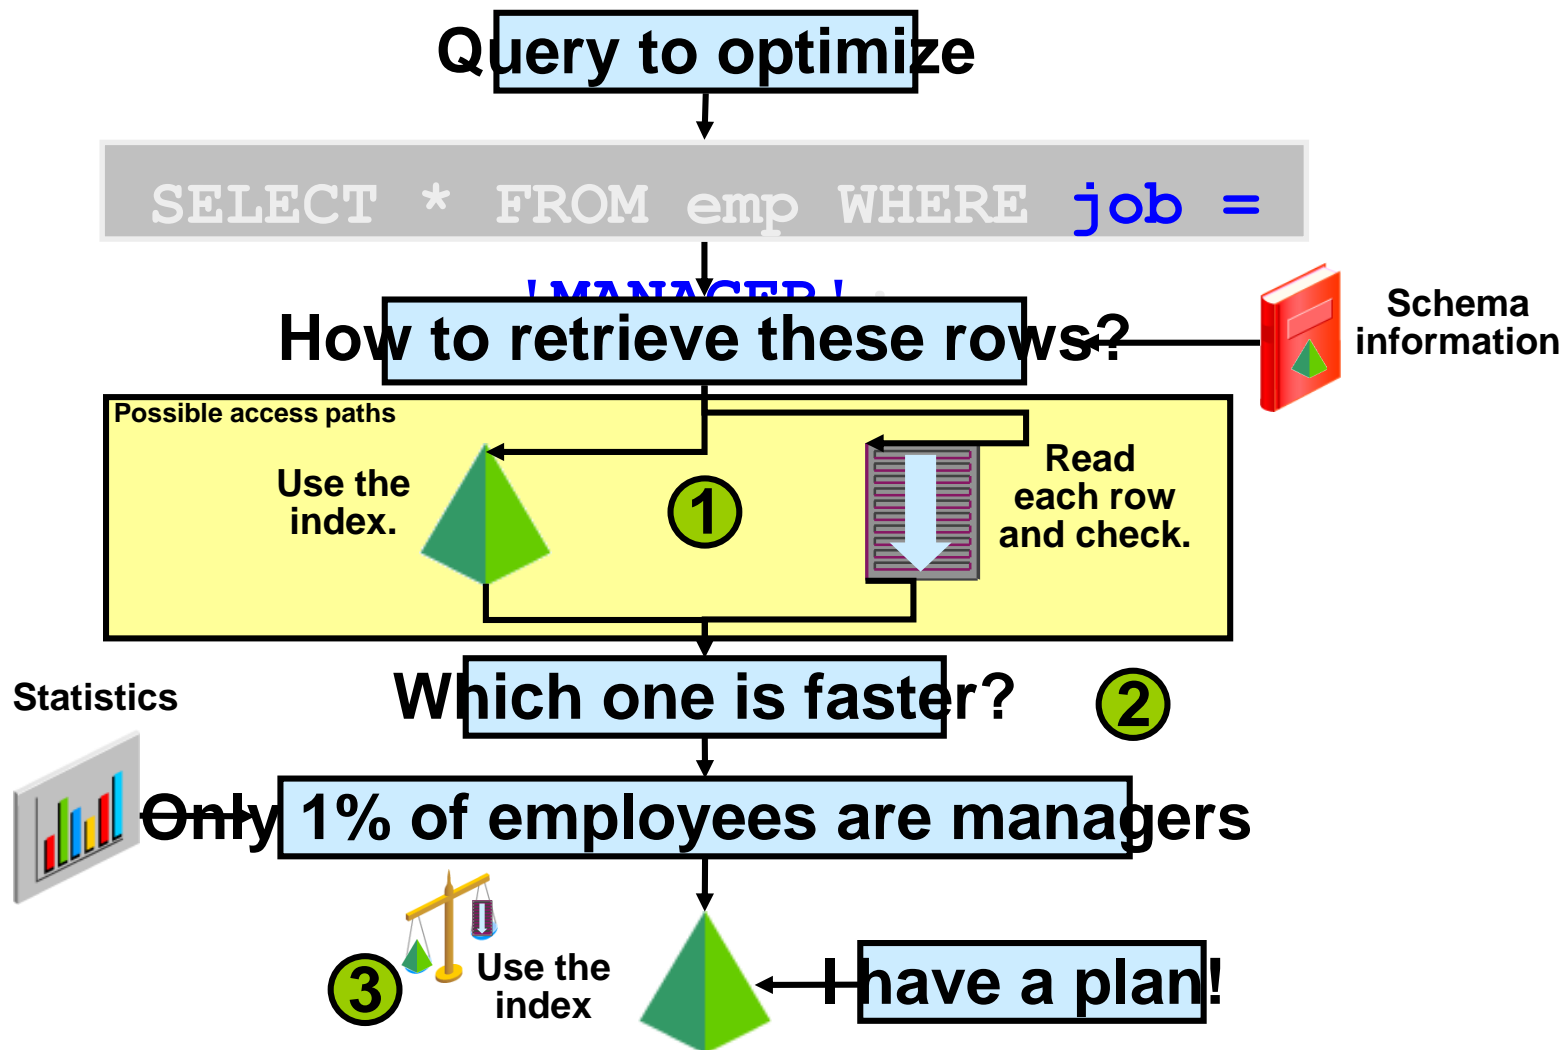
## Writing SQL to Share Cursors

- Create generic code using the following:
  - Stored procedures and packages
  - Database triggers
  - Any other library routines and procedures
- Write to format standards (improves readability):
  - Case
  - White space
  - Comments
  - Object references
  - Bind variables

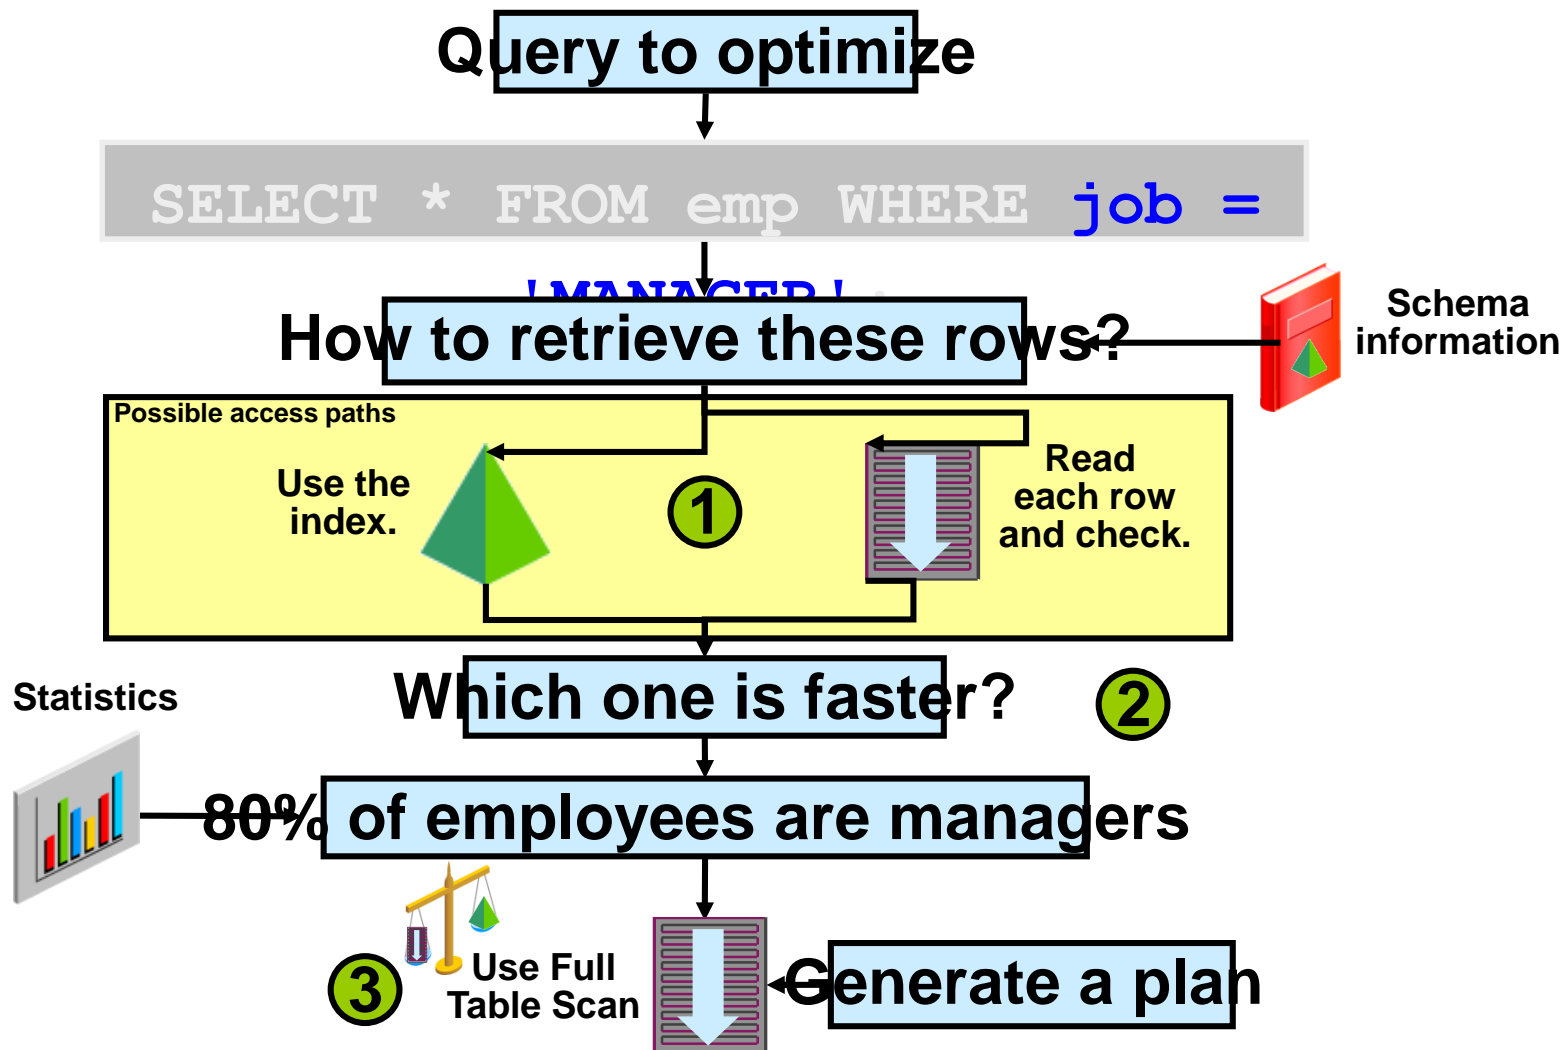


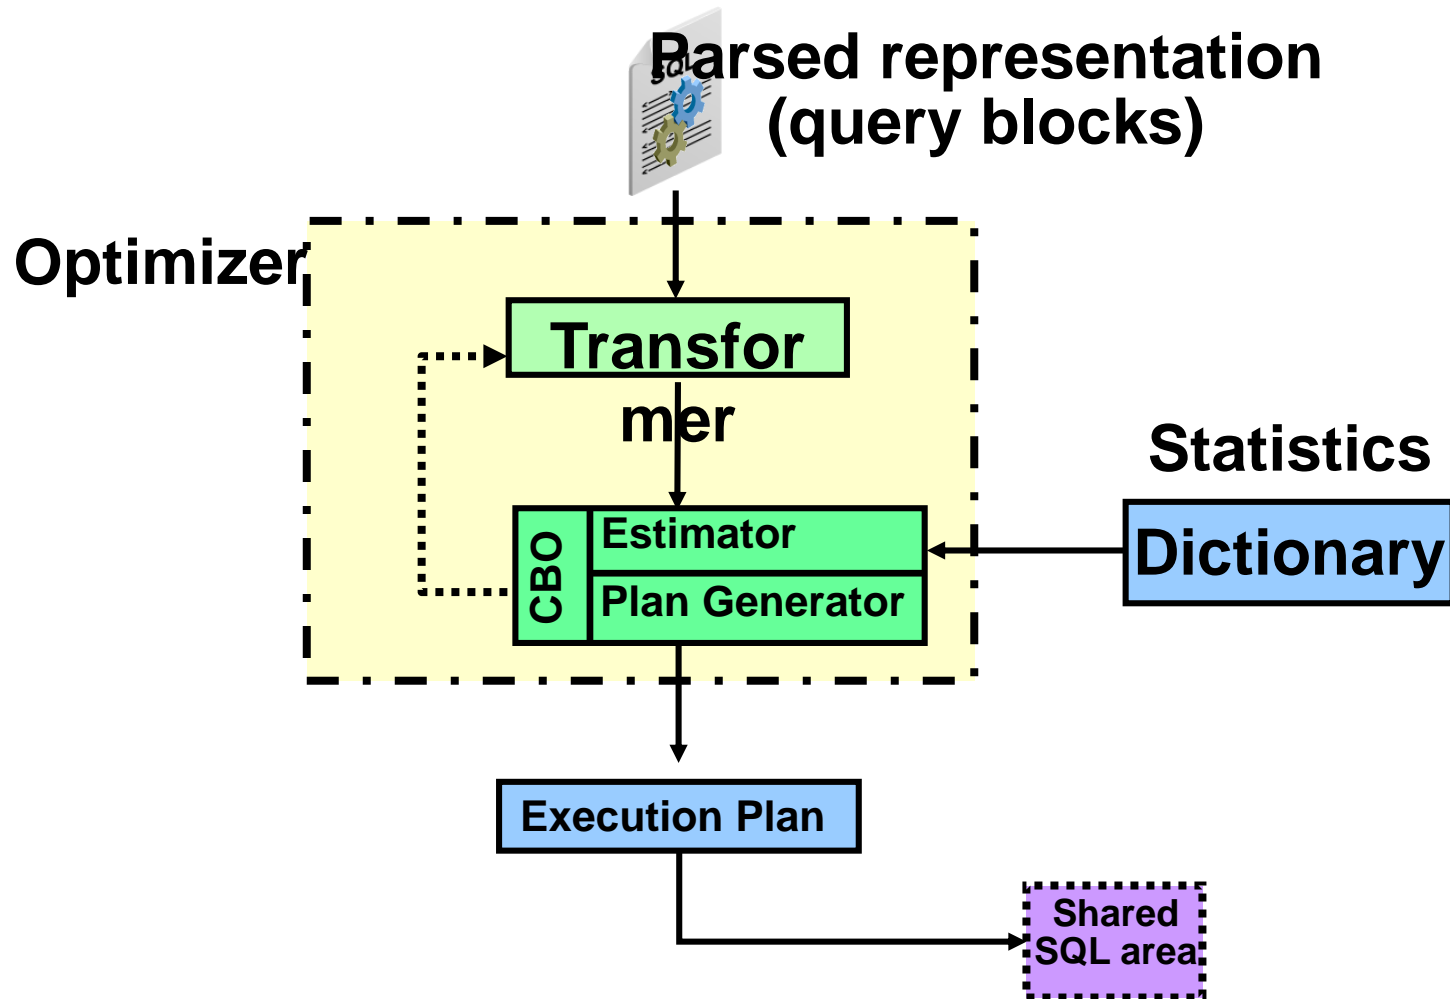
## Performance Checklist

- Set initialization parameters and storage options.
- Verify resource usage of SQL statements.
- Validate connections by middleware.
- Verify cursor sharing.
- Validate migration of all required objects.
- Verify validity and availability of optimizer statistics.











# Cost-Based Optimizer

- Piece of code:
  - Estimator
  - Plan generator
- Estimator determines cost of optimization suggestions made by the plan generator:
  - Cost: Optimizer's best estimate of the number of standardized I/Os made to execute a particular statement optimization
- Plan generator:
  - Tries out different statement optimization techniques
  - Uses the estimator to cost each optimization suggestion
  - Chooses the best optimization suggestion based on cost
  - Generates an execution plan for best optimization



$$\text{Selectivity} = \frac{\text{Number of rows satisfying a condition}}{\text{Total number of rows}}$$

- Selectivity is the estimated proportion of a row set retrieved by a particular predicate or combination of predicates.
- It is expressed as a value between 0.0 and 1.0:
  - High selectivity: Small proportion of rows
  - Low selectivity: Big proportion of rows
- Selectivity computation:
  - If no statistics: Use dynamic sampling
  - If no histograms: Assume even distribution of rows
- Statistic information:
  - `DBA_TABLES` and `DBA_TAB_STATISTICS (NUM_ROWS)`
  - `DBA_TAB_COL_STATISTICS (NUM_DISTINCT, DENSITY, HIGH/LOW_VALUE,...)`



$$\text{Cardinality} = \text{Selectivity} * \text{Total number of rows}$$

- Expected number of rows retrieved by a particular operation in the execution plan
- Vital figure to determine join, filters, and sort costs
- Simple example:
  - The number of distinct values in `DEV_NAME` is 203.
  - The number of rows in `COURSES` (original cardinality) is 1018.
  - Selectivity =  $1/203 = 4.926 \times 10^{-3}$
  - Cardinality =  $(1/203) * 1018 = 5.01$  (rounded off to 6)

```
SELECT days FROM courses WHERE dev name  
= 'ANGEL';
```



```

select e.last_name, d.department_name
      from employees e, departments d
     where e.department_id = d.department_id;

```

```

Join order[1]:  DEPARTMENTS[D]#0  EMPLOYEES[E]#1
               NL Join:  Cost: 41.13  Resp: 41.13  Degree: 1
                   SM cost: 8.01
                   HA cost: 6.51
                   Best:: JoinMethod: Hash
               Cost: 6.51  Degree: 1  Resp: 6.51  Card: 106.00
Join order[2]:  EMPLOYEES[E]#1  DEPARTMENTS[D]#0
               NL Join:  Cost: 121.24  Resp: 121.24  Degree: 1
                   SM cost: 8.01
                   HA cost: 6.51
                   Join order aborted
               Final cost for query block SEL$1 (#0)
               All Rows Plan:
               Best join order: 1

```

Id	Operation	Name	Rows	Bytes	Cost
0	SELECT STATEMENT				7
1	HASH JOIN		106	6042	7
2	TABLE ACCESS FULL	DEPARTMENTS	27	810	3
3	TABLE ACCESS FULL	EMPLOYEES	107	2889	3



## What Is an Execution Plan?

- The execution plan of a SQL statement is composed of small building blocks called row sources for serial execution plans.
- The combination of row sources for a statement is called the execution plan.
- By using parent-child relationships, the execution plan can be displayed in a tree-like structure (text or graphical).





## Where to Find Execution Plans?

- **PLAN\_TABLE (SQL Developer or SQL\*Plus)**
- **V\$SQL\_PLAN (Library Cache)**
- **V\$SQL\_PLAN\_MONITOR (11g)**
- **DBA\_HIST\_SQL\_PLAN (AWR)**
- **STATS\$SQL\_PLAN (Statspack)**
- **SQL management base (SQL plan baselines)**
- **SQL tuning set**
- **Trace files generated by DBMS\_MONITOR**
- **Event 10053 trace file**
- **Process state dump trace file since 10gR2**





## Viewing Execution Plans

- The `EXPLAIN PLAN` command followed by:
  - `SELECT from PLAN_TABLE`
  - `DBMS_XPLAN.DISPLAY()`
- **SQL\*Plus Autotrace:** `SET AUTOTRACE ON`
- `DBMS_XPLAN.DISPLAY_CURSOR()`
- `DBMS_XPLAN.DISPLAY_AWR()`
- `DBMS_XPLAN.DISPLAY_SQLSET()`
- `DBMS_XPLAN.DISPLAY_SQL_PLAN_BASELINE()`



## The `EXPLAIN PLAN` Command

- Generates an optimizer execution plan
- Stores the plan in `PLAN_TABLE`
- Does not execute the statement itself



→→ EXPLAIN PLAN →

[ SET STATEMENT\_ID  
= 'text' ]

→

[ INTO *your plan table* ]

→ FOR *statement* →



```
SQL> EXPLAIN PLAN
      2  SET STATEMENT_ID = 'demo01' FOR
      3  SELECT e.last_name, d.department_name
      4  FROM hr.employees e, hr.departments d
      5  WHERE e.department_id = d.department_id;
```

Explained.

```
SQL>
```

**Note:** The EXPLAIN PLAN command does not actually execute the statement.



## PLAN\_TABLE

- `PLAN_TABLE`:
  - Is automatically created to hold the `EXPLAIN PLAN` output.
  - You can create your own using `utlxplan.sql`.
  - Advantage: SQL is not executed
  - Disadvantage: May not be the actual execution plan
- `PLAN_TABLE` is hierarchical.
- Hierarchy is established with the `ID` and `PARENT_ID` columns.



## Displaying from PLAN\_TABLE: Typical

```
SQL> EXPLAIN PLAN SET STATEMENT_ID = 'demo01' FOR SELECT * FROM emp  
2 WHERE ename = 'KING';
```

Explained.

```
SQL> SET LINESIZE 130  
SQL> SET PAGESIZE 0  
SQL> select * from table(DBMS_XPLAN.DISPLAY());
```

Plan hash value: 3956160932

-----								
Id		Operation		Name		Rows		Bytes   Cost (%CPU)   Time
-----								
0		SELECT STATEMENT				1		37   3 (0)   00:00:01
*		TABLE ACCESS FULL		EMP		1		37   3 (0)   00:00:01
-----								

Predicate Information (identified by operation id):

```
-----  
1 - filter("ENAME"='KING')
```



## Displaying from PLAN\_TABLE: ALL

```
SQL> select * from table(DBMS_XPLAN.DISPLAY(null,null,'ALL'));
```

```
Plan hash value: 3956160932
```

```
-----  
| Id  | Operation                | Name | Rows  | Bytes | Cost (%CPU)| Time     |  
-----  
|  0  | SELECT STATEMENT         |      |    1  |    37 |    3   (0)| 00:00:01 |  
|*  1  |  TABLE ACCESS FULL      | EMP  |    1  |    37 |    3   (0)| 00:00:01 |  
-----
```

```
Query Block Name / Object Alias (identified by operation id):
```

```
-----  
1 - SEL$1 / EMP@SEL$1
```

```
Predicate Information (identified by operation id):
```

```
-----  
1 - filter("ENAME"='KING')
```

```
Column Projection Information (identified by operation id):
```

```
-----  
1 - "EMP"."EMPNO" [NUMBER,22], "EMP"."ENAME" [VARCHAR2,10], "EMP"."JOB" [VARCHAR2,9],  
    "EMP"."MGR" [NUMBER,22], "EMP"."HIREDATE" [DATE,7], "EMP"."SAL" [NUMBER,22],  
    "EMP"."COMM" [NUMBER,22], "EMP"."DEPTNO" [NUMBER,22]
```



```
→→ EXPLAIN PLAN →  
    [ SET STATEMENT_ID  
      = 'text' ]  
  
→ [ INTO your plan table ] →  
  
→ FOR statement →
```



## Displaying from PLAN\_TABLE: ADVANCED



```
select plan_table_output from table(DBMS_XPLAN.DISPLAY(null,null,'ADVANCED
-PROJECTION -PREDICATE -ALIAS'));
```

Plan hash value: 3956160932

Id	Operation	Name	Rows	Bytes	Cost (%CPU)	Time
0	SELECT STATEMENT		1	37	3 (0)	00:00:01
1	TABLE ACCESS FULL	EMP	1	37	3 (0)	00:00:01

Outline Data

/\*+

```
BEGIN_OUTLINE_DATA
FULL(@"SEL$1" "EMP"@"SEL$1")
OUTLINE_LEAF(@"SEL$1")
ALL_ROWS
DB_VERSION('11.1.0.6')
OPTIMIZER_FEATURES_ENABLE('11.1.0.6')
IGNORE_OPTIM_EMBEDDED_HINTS
END_OUTLINE_DATA
```

\*/

## Explain Plan Using SQL Developer



hr x | scott x | sys\_connection x

0.007 seconds

```
select e.email, d.department_name
FROM EMPLOYEES e, departments d
where email like 'A%';
```


Script Output x | Explain Plan x

0.007 seconds

OPERATION	OBJECT_NAME	OPTIONS	COST
SELECT STATEMENT			8
MERGE JOIN		CARTESIAN	8
INDEX	EMP_EMAIL_UK	RANGE SCAN	1
Access Predicates		EMAIL LIKE 'A%'	
Filter Predicates		EMAIL LIKE 'A%'	
BUFFER		SORT	7
TABLE ACCESS	DEPARTMENTS	FULL	2



- Is a SQL\*Plus and SQL Developer facility
- Was introduced with Oracle 7.3
- Needs a `PLAN_TABLE`
- Needs the `PLUSTRACE` role to retrieve statistics from some `V$` views
- By default, produces the execution plan and statistics after running the query
- May not be the execution plan used by the optimizer when using bind peeking (recursive `EXPLAIN PLAN`)





## AUTOTRACE: Examples

- To start tracing statements using `AUTOTRACE`:

- To display:

```
SQL> set autotrace on
```

- To display rows and statistics:

- To get the plan and the statistics only (suppress rows):

```
SQL> set autotrace traceonly explain
```

```
SQL> set autotrace on statistics
```

```
SQL> set autotrace traceonly
```



```
SQL> show autotrace
autotrace OFF
SQL> set autotrace traceonly statistics
SQL> SELECT * FROM oe.products;
```

288 rows selected.

#### Statistics

```
-----
      1334 recursive calls
         0 db block gets
       686 consistent gets
       394 physical reads
         0 redo size
    103919 bytes sent via SQL*Net to client
       629 bytes received via SQL*Net from client
        21 SQL*Net roundtrips to/from client
        22 sorts (memory)
         0 sorts (disk)
       288 rows processed
```



hr x | scott x | sys\_connection x

0.5 seconds

```
select e.email, d.department_name
FROM EMPLOYEES e, departments d
where email like 'A%';
```

Script Output x | Autotrace x

0.5 seconds

OPERATION	OBJECT_NAME	COST	LAST_CR_BUFFER_GETS
SELECT STATEMENT		8	
MERGE JOIN CARTESIAN		8	8
INDEX RANGE SCAN	EMP_EMAIL_UK	1	1
Access Predicates			
EMAIL LIKE 'A%'			
Filter Predicates			

V\$STATNAME Name	V\$MYSTAT Value
recursive calls	1
db block gets	0
consistent gets	8
physical reads	0
redo size	0
bytes sent via SQL*Net to client	1872
bytes received via SQL*Net from client	658
SQL*Net roundtrips to/from client	2
sorts (memory)	2
sorts (disk)	0



## Using the V\$SQL\_PLAN View

- V\$SQL\_PLAN provides a way of examining the execution plan for cursors that are still in the library cache.
- V\$SQL\_PLAN is very similar to PLAN\_TABLE:
  - PLAN\_TABLE shows a theoretical plan that can be used if this statement were to be executed.
  - V\$SQL\_PLAN contains the actual plan used.
- It contains the execution plan of every cursor in the library cache (including child).
- Link to V\$SQL:
  - ADDRESS, HASH\_VALUE, and CHILD\_NUMBER





<b>HASH_VALUE</b>	Hash value of the parent statement in the library cache
<b>ADDRESS</b>	Address of the handle to the parent for this cursor
<b>CHILD_NUMBER</b>	Child cursor number using this execution plan
<b>POSITION</b>	Order of processing for all operations that have the same PARENT_ID
<b>PARENT_ID</b>	ID of the next execution step that operates on the output of the current step
<b>ID</b>	Number assigned to each step in the execution plan
<b>PLAN_HASH_VALUE</b>	Numerical representation of the SQL plan for the cursor

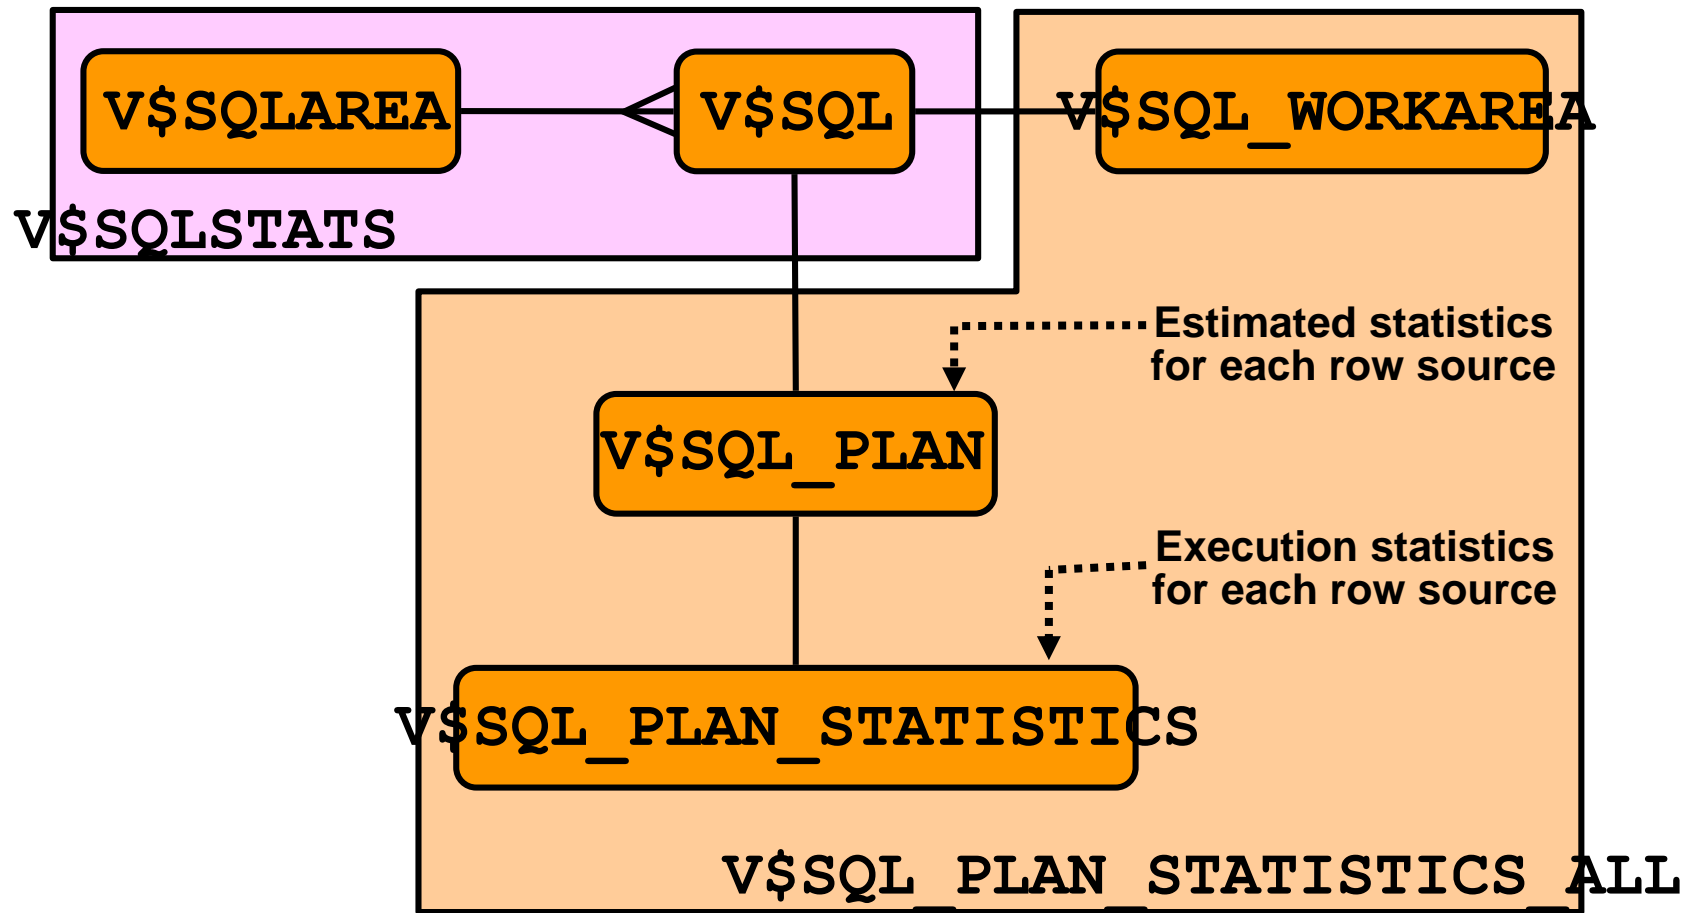
**Note:** This is only a partial listing of the columns.



## The V\$SQL\_PLAN\_STATISTICS View

- V\$SQL\_PLAN\_STATISTICS provides actual execution statistics:
  - STATISTICS\_LEVEL set to ALL
  - The GATHER\_PLAN\_STATISTICS hint
- V\$SQL\_PLAN\_STATISTICS\_ALL enables side-by-side comparisons of the optimizer estimates with the actual execution statistics.

## Links Between Important Dynamic Performance Views





```
SELECT PLAN_TABLE_OUTPUT FROM
TABLE(DBMS_XPLAN.DISPLAY_CURSOR('47ju6102uvq5q'));
```

```
SQL_ID 47ju6102uvq5q, child number 0
```

```
-----
SELECT e.last_name, d.department_name
FROM hr.employees e, hr.departments d WHERE
e.department_id = d.department_id
```

```
Plan hash value: 2933537672
```

Id	Operation	Name	Rows	Bytes	Cost (%CPU)
0	SELECT STATEMENT				6 (100)
1	MERGE JOIN		106	2862	6 (17)
2	TABLE ACCESS BY INDEX ROWID	DEPARTMENTS	27	432	2 (0)
3	INDEX FULL SCAN	DEPT_ID_PK	27		1 (0)
* 4	SORT JOIN		107	1177	4 (25)
5	TABLE ACCESS FULL	EMPLOYEES	107	1177	3 (0)

```
-----
Predicate Information (identified by operation id):
```

```
-----
4 - access("E"."DEPARTMENT_ID"="D"."DEPARTMENT_ID")
    filter("E"."DEPARTMENT_ID"="D"."DEPARTMENT_ID")
```

```
24 rows selected.
```



## Execution Plan Interpretation: Example 1

```
SQL> alter session set statistics_level=ALL;
```

```
Session altered.
```

```
SQL> select /*+ RULE to make sure it reproduces 100% */ ename,job,sal,dname  
from emp,dept where dept.deptno = emp.deptno and not exists (select * from salgrade  
where emp.sal between losal and hisal);
```

```
no rows selected
```

```
SQL> select * from table(dbms_xplan.display_cursor(null,null,'TYPICAL IOSTATS  
LAST'));
```

```
SQL_ID 274019myw3vuf, child number 0
```

```
-----  
...  
Plan hash value: 1175760222
```

-----						
Id	Operation	Name	Starts	A-Rows	Buffers	
-----						
* 1	FILTER		1	0	61	
2	NESTED LOOPS		1	14	25	
3	TABLE ACCESS FULL	EMP	1	14	7	
4	TABLE ACCESS BY INDEX ROWID	DEPT	14	14	18	
* 5	INDEX UNIQUE SCAN	PK_DEPT	14	14	4	
* 6	TABLE ACCESS FULL	SALGRADE	12	12	36	
-----						

```
...
```



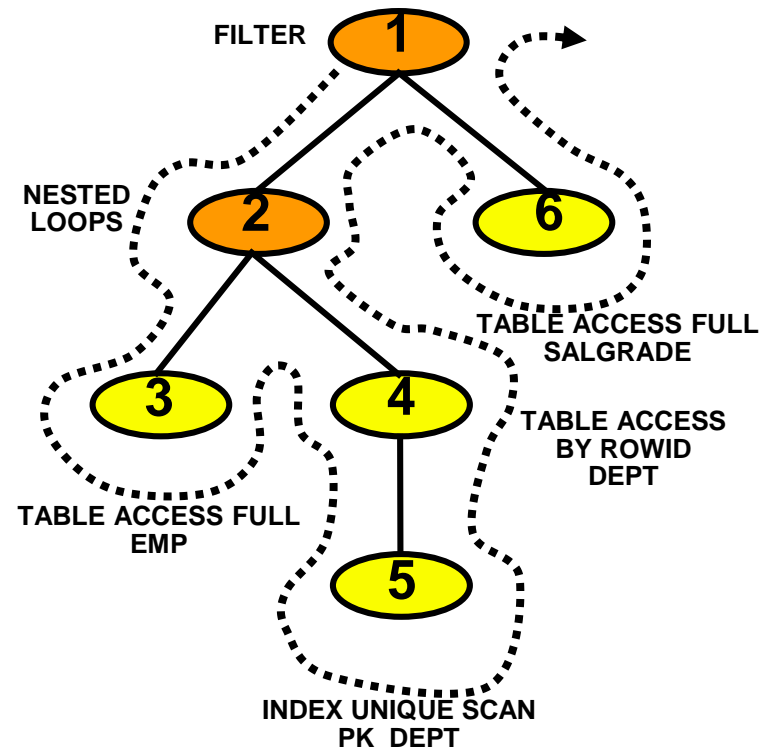
## Execution Plan Interpretation: Example 1

```
SELECT /*+ RULE */ ename,job,sal,dname
FROM emp,dept
WHERE dept.deptno=emp.deptno and not exists(SELECT *
                                           FROM salgrade
                                           WHERE emp.sal between losal and hisal);
```

Id	Operation	Name
0	SELECT STATEMENT	
* 1	FILTER	
2	NESTED LOOPS	
3	TABLE ACCESS FULL	EMP
4	TABLE ACCESS BY INDEX ROWID	DEPT
* 5	INDEX UNIQUE SCAN	PK_DEPT
* 6	TABLE ACCESS FULL	SALGRADE

Predicate Information (identified by operation id):

```
1 - filter( NOT EXISTS
  (SELECT 0 FROM "SALGRADE" "SALGRADE" WHERE
    "HISAL">=:B1 AND "LOSAL"<=:B2))
5 - access("DEPT"."DEPTNO"="EMP"."DEPTNO")
6 - filter("HISAL">=:B1 AND "LOSAL"<=:B2)
```






## Execution Plan Interpretation: Example 2

```
SQL> select /*+ USE_NL(d) use_nl(m) */ m.last_name as dept_manager
2      ,      d.department_name
3      ,      l.street_address
4 from    hr.employees m    join
5          hr.departments d on (d.manager_id = m.employee_id)
6          natural join
7          hr.locations l
8 where   l.city = 'Seattle';
```

```
0  SELECT STATEMENT
1 0  NESTED LOOPS
2 1  NESTED LOOPS
3 2  TABLE ACCESS BY INDEX ROWID LOCATIONS
4 3  INDEX RANGE SCAN LOC_CITY_IX
5 2  TABLE ACCESS BY INDEX ROWID DEPARTMENTS
6 5  INDEX RANGE SCAN DEPT_LOCATION_IX
7 1  TABLE ACCESS BY INDEX ROWID EMPLOYEES
8 7  INDEX UNIQUE SCAN EMP_EMP_ID_PK
```





## Reviewing the Execution Plan

- Drive from the table that has most selective filter.
- Look for the following:
  - Driving table has the best filter
  - Fewest number of rows are returned to the next step
  - The join method is appropriate for the number of rows returned
  - Views are correctly used
  - Unintentional Cartesian products
  - Tables accessed efficiently





- Retrieve all execution plans stored for a particular `SQL_ID`.

```
SQL> SELECT PLAN_TABLE_OUTPUT FROM TABLE (DBMS_XPLAN.DISPLAY_AWR('454rug2yva18w'));
```

```
PLAN_TABLE_OUTPUT
```

```
-----  
SQL_ID 454rug2yva18w  
-----
```

```
select /* example */ * from hr.employees natural join hr.departments
```

```
Plan hash value: 4179021502
```

```
-----  
| Id | Operation | Name | Rows | Bytes | Cost (%CPU) | Time |  
-----  
| 0 | SELECT STATEMENT | | | | 6 (100) | |  
| 1 | HASH JOIN | | 11 | 968 | 6 (17) | 00:00:01 |  
| 2 | TABLE ACCESS FULL | DEPARTMENTS | 11 | 220 | 2 (0) | 00:00:01 |  
| 3 | TABLE ACCESS FULL | EMPLOYEES | 107 | 7276 | 3 (0) | 00:00:01 |  
-----
```

```
SELECT tf.* FROM DBA_HIST_SQLTEXT ht, table  
  (DBMS_XPLAN.DISPLAY_AWR(ht.sql_id,null, null, 'ALL' )) tf  
WHERE ht.sql_text like '%JF%';
```



```
SQL> @$ORACLE_HOME/rdbms/admin/awrsqrpt
```

Specify the Report Type ...

Would you like an HTML report, or a plain text report?

Specify the number of days of snapshots to choose from

Specify the Begin and End Snapshot Ids ...

Specify the SQL Id ...

Enter value for sql\_id: dvza55c7zu0yv

Specify the Report Name ...

## WORKLOAD REPOSITORY SQL Report

### Snapshot Period Summary

DB Name	DB Id	Instance	Inst num	Startup Time	Release	RAC
ORCL	1249102530	orcl	1	14-Jun-10 02:06	11.2.0.1.0	NO
	Snap Id	Snap Time	Sessions	Cursors/Session		
Begin Snap:	218	17-Jun-10 22:00:47	43	6.3		
End Snap:	226	18-Jun-10 04:21:15	40	6.4		
Elapsed:		390.47 (mins)				
DB Time:		5.54 (mins)				

## SQL ID: dvza55c7zu0yv

- 1st Capture and Last Capture Snap IDs refer to Snapshot IDs within the snapshot range
- [SELECT sql\\_id, sql\\_text from DBA\\_HIST\\_SQLTEXT where sql\\_text like '%sa...](#)

#	Plan Hash Value	Total Elapsed Time(ms)	Executions	1st Capture Snap ID	Last Capture Snap ID
1	1258587641	429	1	226	226

[Back to Top](#)

## Plan 1(PHV: 1258587641)

- [Plan Statistics](#)
- [Execution Plan](#)



```

...
select max(cust_credit_limit) from customers where cust_city ='Paris'

```

call	count	cpu	elapsed	disk	query	current	rows
Parse	1	0.00	0.00	0	0	0	0
Execute	1	0.00	0.00	0	0	0	0
Fetch	2	0.00	0.00	1	77	0	1
total	4	0.00	0.00	1	77	0	1

Misses in library cache during parse: 1

Optimizer mode: ALL\_ROWS

Parsing user id: 88

Rows	Row Source Operation
1	SORT AGGREGATE (cr=77 pr=1 pw=0 time=0 us)
77	TABLE ACCESS BY INDEX ROWID CUSTOMERS (cr=77 pr=1 pw=0 time=760 us cost=85 size=1260 card=90)
77	INDEX RANGE SCAN CUST_CUST_CITY_IDX (cr=2 pr=1 pw=0 time=152 us cost=1 size=0 card=90)(object id 78183)



```
tkprof inputfile outputfile [waits=yes|no]
                                [sort=option]
                                [print=n]
                                [aggregate=yes|no]
                                [insert=sqlscriptfile]
                                [sys=yes|no]
                                [table=schema.table]
                                [explain=user/password]
                                [record=statementfile]
                                [width=n]
```



```
tkprof inputfile outputfile [waits=yes|no]
                                [sort=option]
                                [print=n]
                                [aggregate=yes|no]
                                [insert=sqlscriptfile]
                                [sys=yes|no]
                                [table=schema.table]
                                [explain=user/password]
                                [record=statementfile]
                                [width=n]
```



## Output of the `tkprof` Command

➤ There are seven categories of trace statistics:

<b>Count</b>	Number of times the procedure was executed
<b>CPU</b>	Number of seconds to process
<b>Elapsed</b>	Total number of seconds to execute
<b>Disk</b>	Number of physical blocks read
<b>Query</b>	Number of logical buffers read for consistent read
<b>Current</b>	Number of logical buffers read in current mode
<b>Rows</b>	Number of rows processed by the fetch or execute



## Output of the `tkprof` Command

➤ The `tkprof` output also includes the following:

- Recursive SQL statements
- Library cache misses
- Parsing user ID
- Execution plan
- Optimizer mode or hint
- Row source operation

```
...
Misses in library cache during parse: 1
Optimizer mode: ALL_ROWS
Parsing user id: 85

----- Rows      Row Source Operation -----
5  TABLE ACCESS BY INDEX ROWID EMPLOYEES (cr=4 pr=1 pw=0 time=0 us ...
5  INDEX RANGE SCAN EMP_NAME_IX (cr=2 pr=1 pw=0 time=80 us cost=1 ...
...
```



...  
 select max(cust\_credit\_limit) from customers where cust\_city ='Paris'

call	count	cpu	elapsed	disk	query	current	rows
Parse	1	0.00	0.00	0	0	0	0
Execute	1	0.00	0.00	0	0	0	0
Fetch	2	0.02	0.10	72	1459	0	1
total	4	0.02	0.10	72	1459	0	1

Misses in library cache during parse: 1

Optimizer mode: ALL\_ROWS

Parsing user id: 88

Rows	Row Source Operation
1	SORT AGGREGATE (cr=1459 pr=72 pw=0 time=0 us)
77	TABLE ACCESS FULL CUSTOMERS (cr=1459 pr=72 pw=0 time=4104 us cost=405 size=1260 card=90)

...





```

...
select max(cust_credit_limit) from customers where cust_city ='Paris'

```

call	count	cpu	elapsed	disk	query	current	rows
Parse	1	0.00	0.00	0	0	0	0
Execute	1	0.00	0.00	0	0	0	0
Fetch	2	0.00	0.00	1	77	0	1
total	4	0.00	0.00	1	77	0	1

Misses in library cache during parse: 1

Optimizer mode: ALL\_ROWS

Parsing user id: 88

Rows	Row Source Operation
1	SORT AGGREGATE (cr=77 pr=1 pw=0 time=0 us)
77	TABLE ACCESS BY INDEX ROWID CUSTOMERS (cr=77 pr=1 pw=0 time=760 us cost=85 size=1260 card=90)
77	INDEX RANGE SCAN CUST_CUST_CITY_IDX (cr=2 pr=1 pw=0 time=152 us cost=1 size=0 card=90)(object id 78183)



```

...
select max(cust_credit_limit) from customers where cust_city ='Paris'

```

call	count	cpu	elapsed	disk	query	current	rows
Parse	1	0.00	0.00	0	0	0	0
Execute	1	0.00	0.00	0	0	0	0
Fetch	2	0.00	0.00	1	77	0	1
total	4	0.00	0.00	1	77	0	1

Misses in library cache during parse: 1

Optimizer mode: ALL\_ROWS

Parsing user id: 88

Rows	Row Source Operation
1	SORT AGGREGATE (cr=77 pr=1 pw=0 time=0 us)
77	TABLE ACCESS BY INDEX ROWID CUSTOMERS (cr=77 pr=1 pw=0 time=760 us cost=85 size=1260 card=90)
77	INDEX RANGE SCAN CUST_CUST_CITY_IDX (cr=2 pr=1 pw=0 time=152 us cost=1 size=0 card=90)(object id 78183)



- For all sessions in the database:

```
EXEC dbms_monitor.DATABASE_TRACE_ENABLE(TRUE,TRUE);
```

```
EXEC dbms_monitor.DATABASE_TRACE_DISABLE();
```

```
EXEC dbms_monitor.SESSION_TRACE_ENABLE(session_id=>27, serial_num=>60, waits=>TRUE, binds=>FALSE);
```

```
EXEC dbms_monitor.SESSION_TRACE_DISABLE(session_id=>27, serial_num=>60);
```



- Enabling trace:

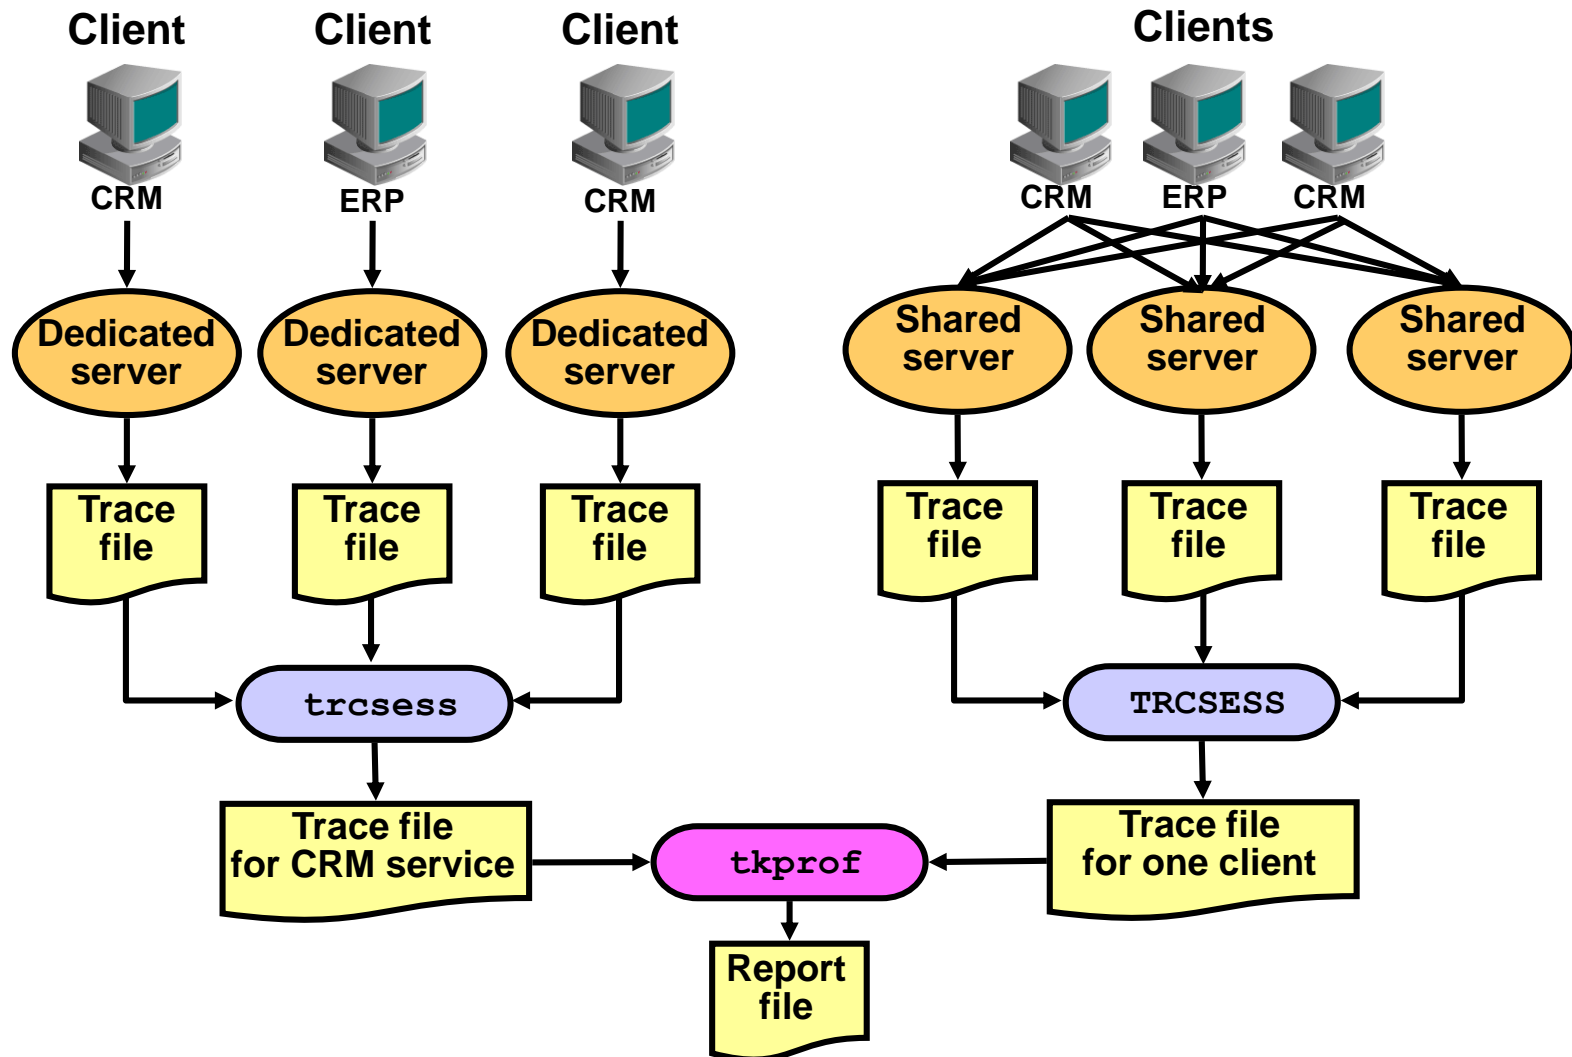
```
EXEC DBMS_SESSION.SESSION_TRACE_ENABLE(waits =>  
TRUE, binds => FALSE);
```

- Disabling trace:

```
EXEC DBMS_SESSION.SESSION_TRACE_DISABLE();
```


- Easily identifying your trace files:

```
alter session set  
tracefile_identifier='mytraceid';
```





```
trcse  [output=output_file_name]
       [session=session_id]
       [clid=client_identifier]
       [service=service_name]
       [action=action_name]
       [module=module_name]
       [<trace file names>]
```





```
exec dbms_session.set_identifier('HR session');
```

First session

Second session

```
exec dbms_session.set_identifier('HR session');
```

```
exec DBMS_MONITOR.CLIENT_ID_TRACE_ENABLE( -  
  client_id=>'HR session', waits => FALSE, -  
  binds => FALSE);
```

Third session

```
select * from employees;  
...
```

```
select * from departments;  
...
```

```
exec DBMS_MONITOR.CLIENT_ID_TRACE DISABLE( -  
  client_id => 'HR session');
```

```
trcsess output=mytrace.trc clientid='HR session'  
$ORACLE_BASE/diag/rdbms/orcl/orcl/trace/*.trc
```



## SQL Trace File Contents

- Parse, execute, and fetch counts
- CPU and elapsed times
- Physical reads and logical reads
- Number of rows processed
- Misses on the library cache
- Username under which each parse occurred
- Each commit and rollback
- Wait event and bind data for each SQL statement
- Row operations showing the actual execution plan of each SQL statement
- Number of consistent reads, physical reads, physical writes, and time elapsed for each operation on a row



## SQL Trace File Contents: Example




```
*** [ Unix process pid: 15911 ]
*** 2010-07-29 13:43:11.327
*** 2010-07-29 13:43:11.327
*** 2010-07-29 13:43:11.327
*** 2010-07-29 13:43:11.327
...
=====
PARSING IN CURSOR #2 len=23 dep=0 uid=85 oct=3 lid=85 tim=1280410994003145 hv=40
69246757 ad='4cd57ac0' sqlid='f34thrbt8rjt5'
      select * from employees
      END OF STMT
PARSE #2:c=3000,e=2264,p=0,cr=0,cu=0,mis=1,r=0,dep=0,og=1,plh=1445457117,
      tim=1280410994003139
EXEC #2:c=0,e=36,p=0,cr=0,cu=0,mis=0,r=0,dep=0,og=1,plh=1445457117,
      tim=1280410994003312
FETCH #2:c=0,e=215,p=0,cr=3,cu=0,mis=0,r=1,dep=0,og=1,plh=1445457117,
      tim=1280410994003628
FETCH #2:c=0,e=89,p=0,cr=5,cu=0,mis=0,r=15,dep=0,og=1,plh=1445457117,
      tim=1280410994004232
...
FETCH #2:c=0,e=60,p=0,cr=1,cu=0,mis=0,r=1,dep=0,og=1,plh=1445457117,
      tim=1280410994107857
STAT #2 id=1 cnt=107 pid=0 pos=1 obj=73933 op='TABLE ACCESS FULL EMPLOYEES (cr=15
pr=0 pw=0 time=0 us cost=3 size=7383 card=107)'
XCTEND rlbk=0, rd_only=1, tim=1280410994108875
=====
```

## Formatting SQL Trace Files: Overview



➤ Use the `tkprof` utility to format your SQL trace files:

- Sort raw trace file to exhibit top SQL statements
- Filter dictionary statements





### Structures

#### Tables

### Access Paths

Full Table Scan

Rowid Scan

Sample Table Scan

#### Indexes

Index Scan (Unique)

Index Scan (Range)

Index Scan (Full)

Index Scan (Fast Full)

Index Scan (Skip)

Index Scan (Index Join)


Using Bitmap Indexes

Combining Bitmap Indexes

# Full Table Scan



- Performs multiblock reads  
(here `DB_FILE_MULTIBLOCK_READ_COUNT = 4`)
- Reads all formatted blocks below the high-water mark
- May filter rows
- Is faster than index  
range scans for large amount of data



0.38699999 seconds | scott

```
select * from emp where ename='King';
```

Explain Plan x


0.387 seconds

OPERATION	OBJECT_NAME	OPTIONS
SELECT STATEMENT		
TABLE ACCESS	EMP	<b>FULL</b>
Filter Predicates ENAME='King'		



```
select * from scott.emp where rowid='AAQ+LAAEAAAAfAAJ';
```


Id	Operation	Name	Rows	Bytes	Cost
0	SELECT STATEMENT		1	37	1
1	TABLE ACCESS BY USER ROWID	EMP	1	37	1





```
SELECT * FROM emp SAMPLE BLOCK (10) SEED (1);
```

Id	Operation	Name	Rows	Bytes	Cost (%CPU)
0	SELECT STATEMENT		4	99	2 (0)
1	TABLE ACCESS SAMPLE	EMP	4	99	2 (0)





## Automatic Workload Repository (AWR)

- Collects, processes, and maintains performance statistics for problem-detection and self-tuning purposes
- Statistics include:
  - Object statistics
  - Time-model statistics
  - Some system and session statistics
  - Active Session History (ASH) statistics
- Automatically generates snapshots of the performance data



# Indexes: Overview

## ➤ Indexes

- Storage techniques:
  - B\*-tree indexes: The default and the most common
    - Normal
    - Function based: Precomputed value of a function or expression
    - Index-organized table (IOT)
  - Bitmap indexes
  - Cluster indexes: Defined specifically for cluster
- Index attributes:
  - Key compression
  - Reverse key
  - Ascending, descending
- Domain indexes: Specific to an application or cartridge






Table data retrieved by using rowid


# Index Scans



## Types of index scans:

- Unique
- Min/Max
- Range (Descending)
- Skip
- Full and fast full
- Index join


B-Tree index IX\_EMP



## Index Unique Scan



index UNIQUE Scan PK\_EMP



```
create unique index PK_EMP on EMP(empno)
```

```
select * from emp where empno = 9999;
```

Explain Plan x				
📌   0.024 seconds				
OPERATION	OBJECT_NAME	OPTIONS	COST	CARDINALITY
SELECT STATEMENT			1	1
TABLE ACCESS	EMP	BY INDEX ROWID	1	1
INDEX	PK_EMP	UNIQUE SCAN	0	1
Access Predicat				
EMPNO=7839				


# Index Range Scan



Index Range SCAN I\_DEPTNO


```
create index I_DEPTNO on EMP (deptno);

select /*+ INDEX(EMP I_DEPTNO) */ *
from emp where deptno = 10 and sal > 1000;
```



Script Output x Explain Plan x				
0.012 seconds				
OPERATION	OBJECT_NAME	OPTIONS	COST	CARDINALITY
SELECT STATEMENT			2	4
TABLE ACCESS	EMP	BY INDEX ROWID	2	4
Filter Predicates				
SAL>1000				
INDEX	I_DEPTNO	RANGE SCAN	1	5
Access Predicat				
DEPTNO=10				

## Index Range Scan: Descending




```
select * from emp where deptno>20 order by deptno desc;
```

Script Output x Explain Plan x				
📌   0.426 seconds				
OPERATION	OBJECT_NAME	OPTIONS	COST	CARDINALITY
SELECT STATEMENT			2	7
TABLE ACCESS	EMP	BY INDEX ROWID	2	7
INDEX	I_DEPTNO	RANGE SCAN DESCENDING	1	7
Access Predicat				
DEPTNO>20				

# Descending Index Range Scan



Index Range SCAN IX\_D




```
drop index I_Deptno;  
create index IX_D on EMP(deptno);  
select * from emp where deptno < 30;
```

Script Output x Explain Plan x				
0.011 seconds				
OPERATION	OBJECT_NAME	OPTIONS	COST	CARDINALITY
SELECT STATEMENT			2	9
TABLE ACCESS	EMP	BY INDEX ROWID	2	9
INDEX	IX_D	RANGE SCAN	1	1
Access Predicates				
SYS_OP_DESCEND(DEPTNO)>HEXTORAW('3EE0FF')				
Filter Predicates				
SYS_OP_UNDESCEND(SYS_OP_DESCEND(DEPTNO))<30				

## Index Range Scan: Function-Based



Index Range SCAN IX\_FBI



```
create index IX_FBI on EMP(UPPER(ENAME))
```

```
select * from emp where upper(ENAME) like 'A%';
```


OPERATION	OBJECT_NAME	OPTIONS	COST	CARDINALITY
SELECT STATEMENT			2	1
TABLE ACCESS	EMP	BY INDEX RO...	2	1
INDEX	IX_FBI	RANGE SCAN	1	1
Access Predicates				
UPPER(ENAME) LIKE 'A%'				
Filter Predicates				
UPPER(ENAME) LIKE 'A%'				

## Index Full Scan



```
create index I_DEPTNO on EMP(deptno
```

```
select *  
from emp  
where sal > 1000 and deptno  
order by deptno;
```



Script Output x Explain Plan x				
0.012 seconds				
OPERATION	OBJECT_NAME	OPTIONS	COST	CARDINALITY
SELECT STATEMENT			2	13
TABLE ACCESS	EMP	BY INDEX RO...	2	13
Filter Predicates SAL>1000				
INDEX	I_DEPTNO	FULL SCAN	1	14
Filter Predicates DEPTNO IS NOT NULL				



# Index Fast Full Scan



db file multiblock read c

ount = 4  
multiblock read      multiblock read


## LEGEND:

SH=segment header

R=root block

B=branch block

L=leaf block



```
select /*+ INDEX_FFS(EMP I_DEPTNO) */ deptno from emp
       where deptno is not null;
```

Script Output x Explain Plan x				
0.01 seconds				
OPERATION	OBJECT_NAME	OPTIONS	COST	CARDINALITY
SELECT STATEMENT			2	14
INDEX	I_DEPTNO	FAST FULL SCAN	2	14
Filter Predicates				
DEPTNO IS NOT NULL				




Table data retrieved by using rowid


# Index Scans



## Types of index scans:

- Unique
- Min/Max
- Range (Descending)
- Skip
- Full and fast full
- Index join


B-Tree index IX\_EMP



## Index Unique Scan



index UNIQUE Scan PK\_EMP



```
create unique index PK_EMP on EMP(empno)
```


```
select * from emp where empno = 9999;
```

Explain Plan x				
📌   0.024 seconds				
OPERATION	OBJECT_NAME	OPTIONS	COST	CARDINALITY
SELECT STATEMENT			1	1
TABLE ACCESS	EMP	BY INDEX ROWID	1	1
INDEX	PK_EMP	UNIQUE SCAN	0	1
Access Predicat				
EMPNO=7839				

# Index Range Scan



Index Range SCAN I\_DEPTNO



```
create index I_DEPTNO on EMP (deptno);


select /*+ INDEX(EMP I_DEPTNO) */ *
from emp where deptno = 10 and sal > 1000;
```

Script Output x Explain Plan x

0.012 seconds

OPERATION	OBJECT_NAME	OPTIONS	COST	CARDINALITY
SELECT STATEMENT			2	4
TABLE ACCESS	EMP	BY INDEX ROWID	2	4
Filter Predicates				
SAL>1000				
INDEX	I_DEPTNO	RANGE SCAN	1	5
Access Predicat				
DEPTNO=10				

## Index Range Scan: Descending




```
select * from emp where deptno>20 order by deptno desc;
```

Script Output x Explain Plan x				
📌   0.426 seconds				
OPERATION	OBJECT_NAME	OPTIONS	COST	CARDINALITY
SELECT STATEMENT			2	7
TABLE ACCESS	EMP	BY INDEX ROWID	2	7
INDEX	I_DEPTNO	RANGE SCAN DESCENDING	1	7
Access Predicat				
DEPTNO>20				

## Descending Index Range Scan



Index Range SCAN IX\_D




```
drop index I_Deptno;  
create index IX_D on EMP(deptno);  
select * from emp where deptno < 30;
```

Script Output x Explain Plan x				
0.011 seconds				
OPERATION	OBJECT_NAME	OPTIONS	COST	CARDINALITY
SELECT STATEMENT			2	9
TABLE ACCESS	EMP	BY INDEX ROWID	2	9
INDEX	IX_D	RANGE SCAN	1	1
Access Predicates				
SYS_OP_DESCEND(DEPTNO)>HEXTORAW('3EE0FF')				
Filter Predicates				
SYS_OP_UNDESCEND(SYS_OP_DESCEND(DEPTNO))<30				

## Index Range Scan: Function-Based



Index Range SCAN IX\_FBI



```
create index IX_FBI on EMP(UPPER(ENAME))
```

```
select * from emp where upper(ENAME) like 'A%';
```

OPERATION	OBJECT_NAME	OPTIONS	COST	CARDINALITY
SELECT STATEMENT			2	1
TABLE ACCESS	EMP	BY INDEX RO...	2	1
INDEX	IX_FBI	RANGE SCAN	1	1
Access Predicates				
UPPER(ENAME) LIKE 'A%'				
Filter Predicates				
UPPER(ENAME) LIKE 'A%'				




# Index Full Scan



```
create index I_DEPTNO on EMP(deptno)

select *
from emp
where sal > 1000 and deptno is not null
order by deptno;
```



Script Output x Explain Plan x				
0.012 seconds				
OPERATION	OBJECT_NAME	OPTIONS	COST	CARDINALITY
SELECT STATEMENT			2	13
TABLE ACCESS	EMP	BY INDEX RO...	2	13
Filter Predicates SAL>1000				
INDEX	I_DEPTNO	FULL SCAN	1	14
Filter Predicates DEPTNO IS NOT NULL				

# Index Fast Full Scan



db file multiblock read c

ount = 4  
multiblock read      multiblock read


## LEGEND:

SH=segment header

R=root block

B=branch block

L=leaf block



```
select /*+ INDEX_FFS(EMP I_DEPTNO) */ deptno from emp
       where deptno is not null;
```


Script Output x Explain Plan x				
0.01 seconds				
OPERATION	OBJECT_NAME	OPTIONS	COST	CARDINALITY
SELECT STATEMENT			2	14
INDEX	I_DEPTNO	FAST FULL SCAN	2	14
Filter Predicates				
DEPTNO IS NOT NULL				



## *Index on (DEPTNO, SAL)*

```
create index IX_SS on EMP(DEPTNO,SAL);
```

```
select /*+ index_ss(EMP IX_SS) */ * from emp where SAL < 1500;
```

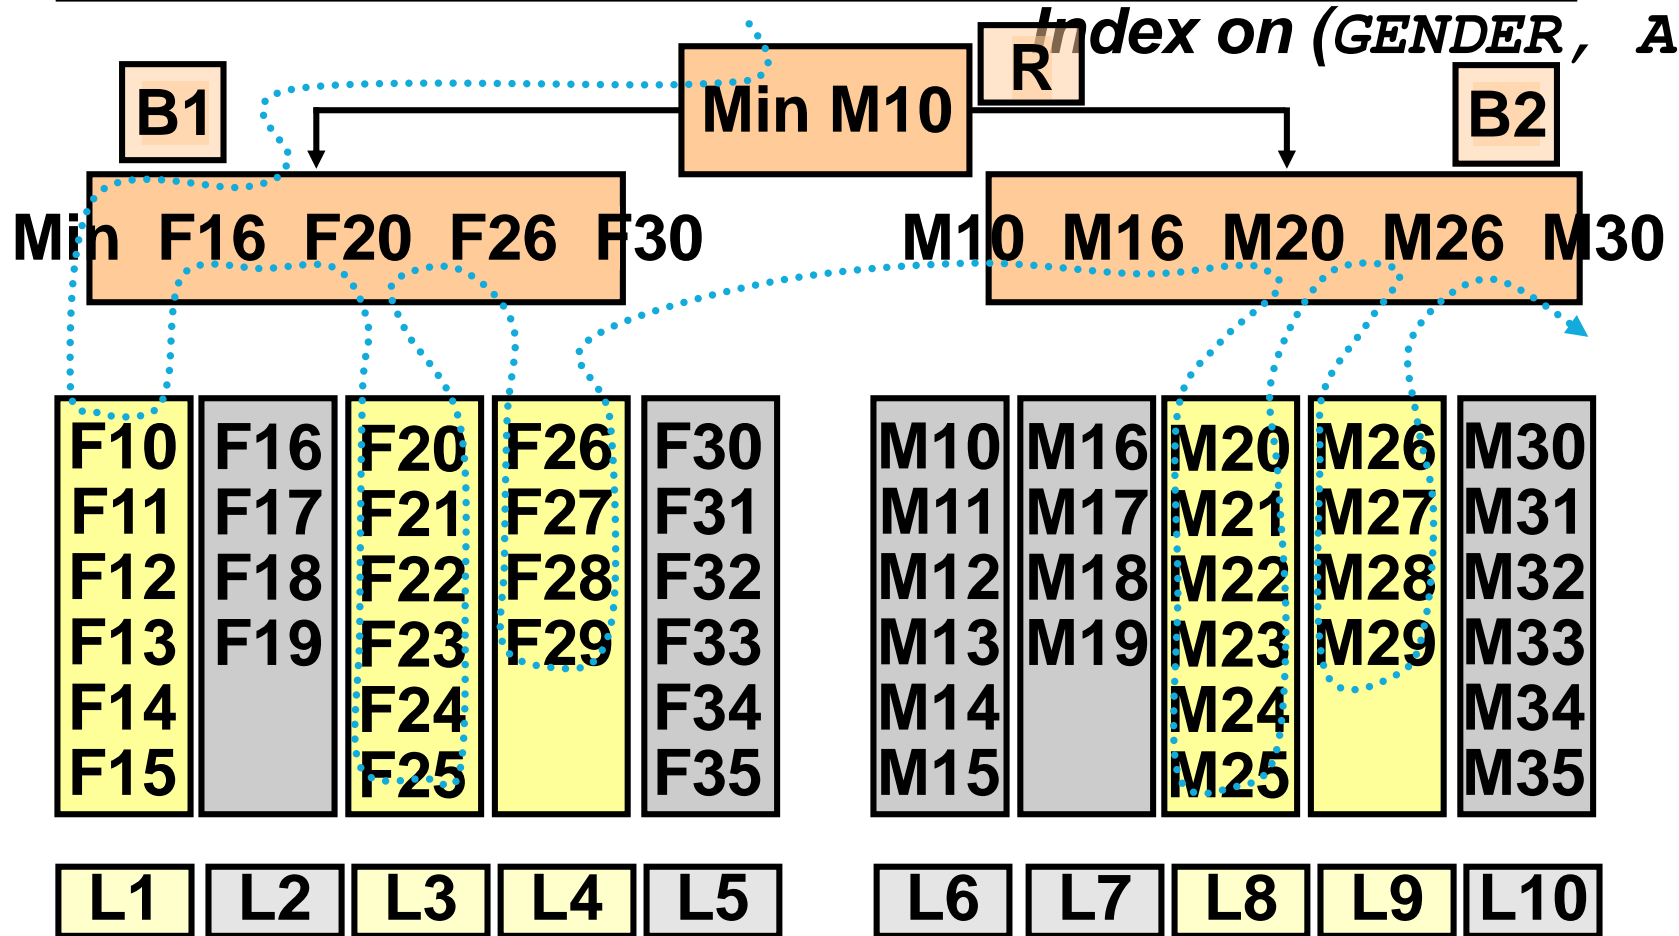


Script Output x Explain Plan x				
0.01 seconds				
OPERATION	OBJECT_NAME	OPTIONS	COST	CARDINALITY
SELECT STATEMENT			4	2
TABLE ACCESS	EMP	BY INDEX ROWID	4	2
INDEX	IX_SS	SKIP SCAN	3	2
Access Predicates				
SAL<1500				
Filter Predicates				
SAL<1500				



SELECT \* FROM employees WHERE age BETWEEN 20 AND 29

Index on (GENDER, AGE)



## Index Join Scan



```
alter table emp modify (SAL not null, ENAME not null);  
create index I_ENAME on EMP(ename);  
create index I_SAL on EMP(sal);
```

```
select /*+ INDEX_JOIN(e) */ ename, sal from emp e;
```

Statement Output x

Autotrace x

1.563 seconds

OPERATION	OBJECT_NAME	COST
SELECT STATEMENT		3
VIEW	index\$_join\$_001	3
type="db_version"		
11.2.0.1		
HASH JOIN		
Access Predicates		
ROWID=ROWID		
INDEX FAST FULL SCAN	I_ENAME	1
INDEX FAST FULL SCAN	I_SAL	1

## B\*-tree Indexes and Nulls



```
create table nulltest ( col1 number, col2 number not null);  
create index nullind1 on nulltest (col1);  
create index notnullind2 on nulltest (col2);
```

```
select /*+ index(t nullind1) */ col1 from nulltest t;
```

Script Output x Explain Plan x

0.864 seconds

OPERATION	OBJECT_NAME	OPTIONS	COST	CARDINALITY
SELECT STATEMENT			2	1
TABLE ACCESS	NULLTEST	FULL	2	1

```
select col1 from nulltest t where col1=10;
```

Script Output x Explain Plan x

1.105 seconds

OPERATION	OBJECT_NAME	OPTIONS	COST	CARDINALITY
SELECT STATEMENT			1	1
INDEX	NULLIND1	RANGE SCAN	1	1
Access Predicates				
COL1=10				

```
select /*+ index(t notnullind2) */ col2 from nulltest t;
```

Script Output x Explain Plan x

0.034 seconds

OPERATION	OBJECT_NAME	OPTIONS	COST	CARDINALITY
SELECT STATEMENT			1	1
INDEX	NOTNULLIND2	FULL SCAN	1	1

## Using Indexes: Considering Nullable Columns



Column Null?

SSN	Y
FNAME	Y
LNAME	N
•	
•	
•	

PERSON

```
CREATE UNIQUE INDEX person_ssn_ix
ON person(ssn);
```

```
SELECT COUNT(*) FROM person;
```

```
SELECT STATEMENT      |
  SORT AGGREGATE      |
TABLE ACCESS FULL | PERSON
```

```
DROP INDEX person_ssn_ix;
```

Column Null?



SSN	N
FNAME	Y
LNAME	N
•	
•	
•	

PERSON

```
ALTER TABLE person ADD CONSTRAINT pk_ssn
PRIMARY KEY (ssn);
```

```
SELECT /*+ INDEX(person) */ COUNT(*) FROM
      person;
```

```
SELECT STATEMENT      |
  SORT AGGREGATE      |
INDEX FAST FULL SCAN | PK_SSN
```





## Index-Organized Table Scans







```
create table iotemp
( empno number(4) primary key, ename varchar2(10) not null,
  job varchar2(9), mgr number(4), hiredate date,
  sal number(7,2) not null, comm number(7,2), deptno number(2))
organization index;
```

```
select * from iotemp where empno=9999;
```

Script Output x

Explain Plan x


 | 0.734 seconds

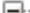


OPERATION	OBJECT_NAME	OPTIONS	COST	CARDINALITY
 SELECT STATEMENT			1	
 INDEX	SYS_IOT_TOP_7..	UNIQUE SCAN	1	
 Access Predicates EMPNO=9999				

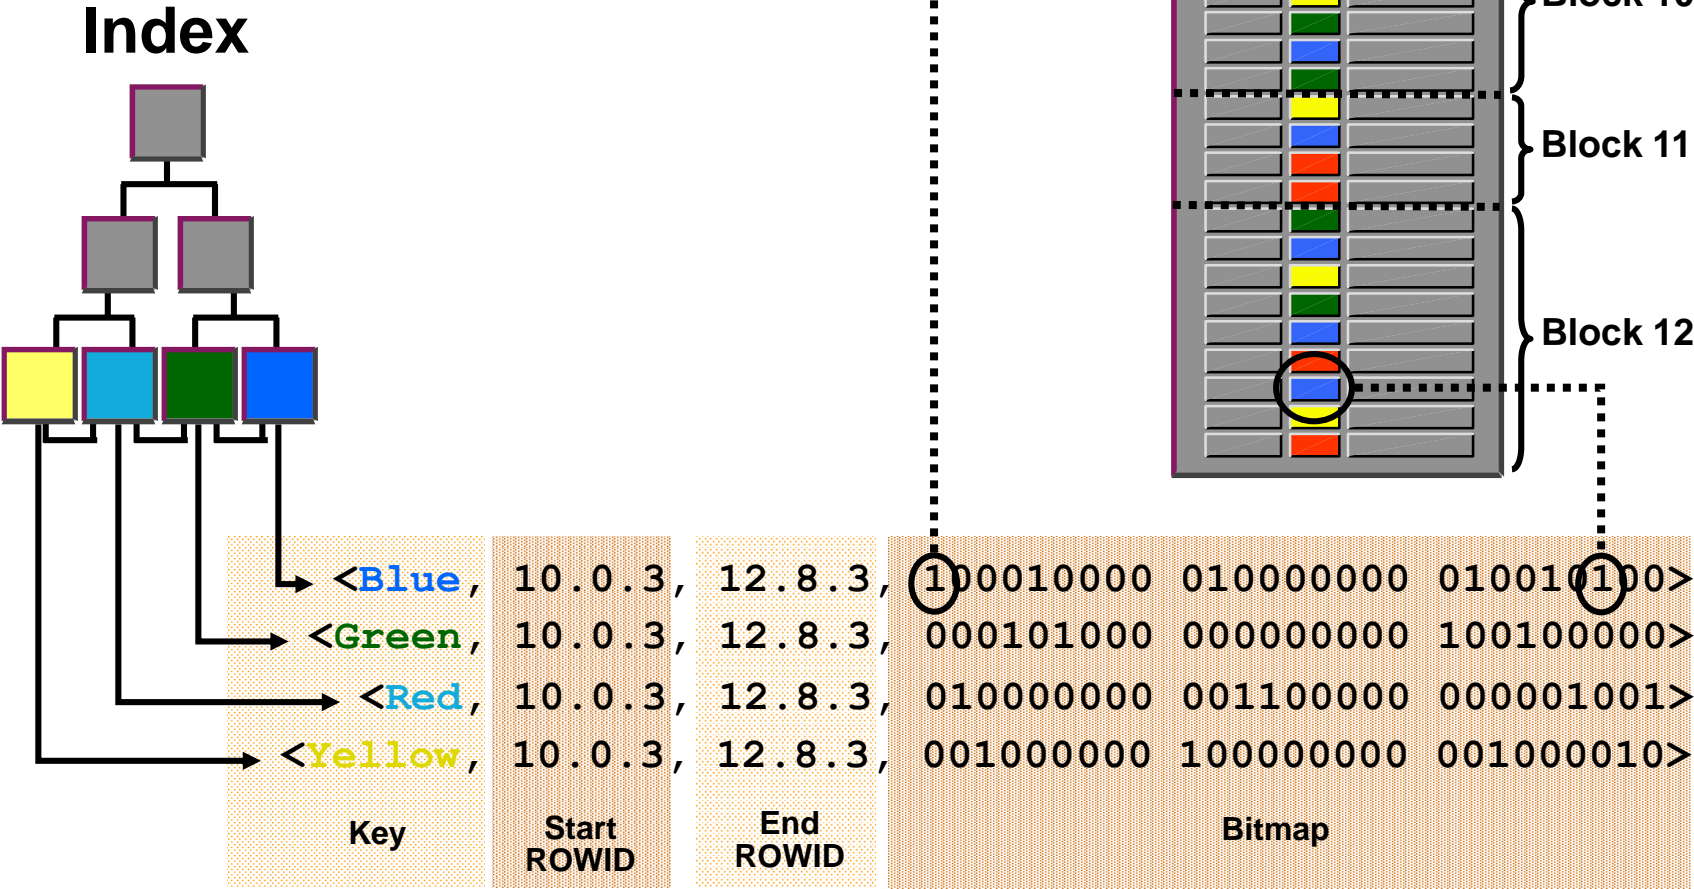
```
select * from iotemp where sal>1000;
```

Script Output x

Explain Plan x

 | 0.007 seconds

OPERATION	OBJECT_NAME	OPTIONS	COST	CARDINALITY
 SELECT STATEMENT			2	1
 INDEX	SYS_IOT_TOP_7...	FAST FULL SCAN	2	1
 Filter Predicates				
SAL>1000				



## Bitmap Index Access: Examples



```
SELECT * FROM PERF_TEAM WHERE country='FR';
```

Id	Operation	Name	Rows	Bytes
0	SELECT STATEMENT		1	45
1	TABLE ACCESS BY INDEX ROWID	PERF_TEAM	1	45
2	BITMAP CONVERSION TO ROWIDS			
3	BITMAP INDEX SINGLE VALUE	IX_B2		

```
Predicate: 3 - access("COUNTRY"='FR')
```

```
SELECT * FROM PERF_TEAM WHERE country>'FR';
```



Id	Operation	Name	Rows	Bytes
0	SELECT STATEMENT		1	45
1	TABLE ACCESS BY INDEX ROWID	PERF_TEAM	1	45
2	BITMAP CONVERSION TO ROWIDS			
3	BITMAP INDEX RANGE SCAN	IX_B2		

```
Predicate: 3 - access("COUNTRY">'FR') filter("COUNTRY">'FR')
```

## Combining Bitmap Indexes: Examples



```
SELECT * FROM PERF_TEAM WHERE country in('FR','DE');
```



```
SELECT * FROM EMEA_PERF_TEAM T WHERE country='FR' and gender='M';
```

## Combining Bitmap Index Access Paths



```
SELECT * FROM PERF_TEAM WHERE country in ('FR','DE');
```

Id	Operation	Name	Rows	Bytes
0	SELECT STATEMENT		1	45
1	INLIST ITERATOR			
2	TABLE ACCESS BY INDEX ROWID	PERF_TEAM	1	45
3	BITMAP CONVERSION TO ROWIDS			
4	BITMAP INDEX SINGLE VALUE	IX_B2		

Predicate: 4 - access("COUNTRY"='DE' OR "COUNTRY"='FR')

```
SELECT * FROM PERF_TEAM WHERE country='FR' and gender='M';
```

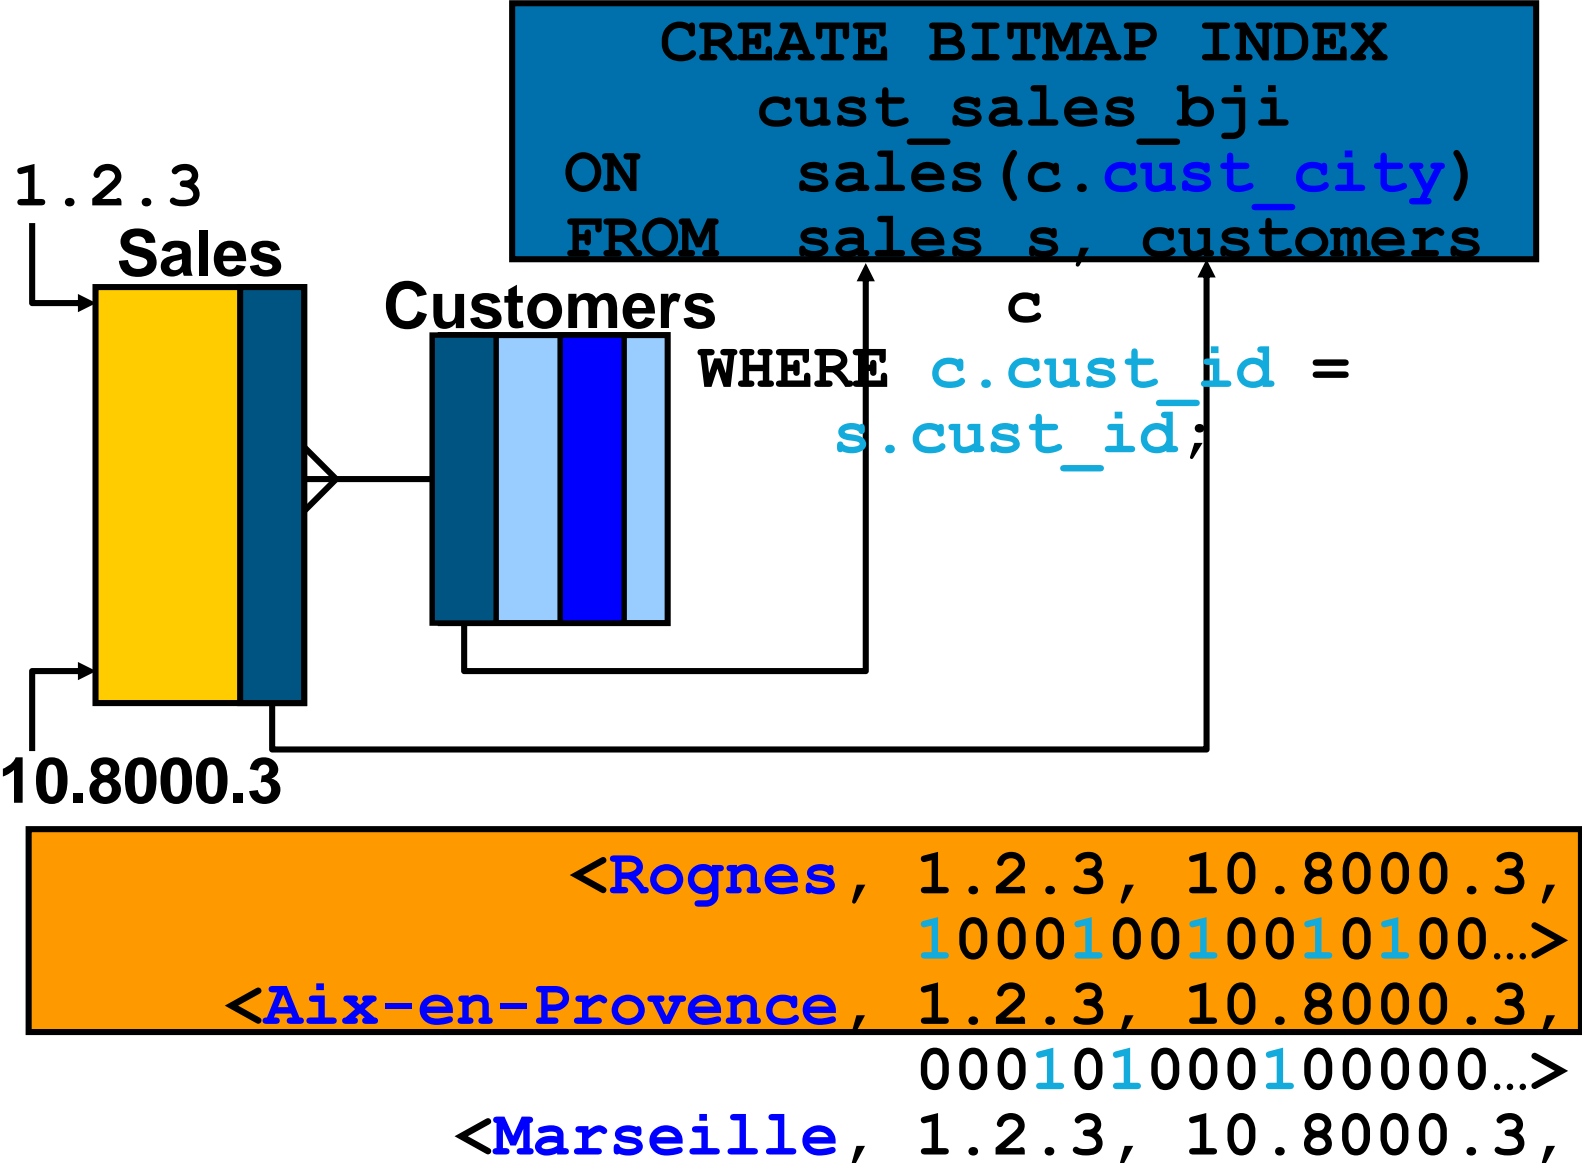
Id	Operation	Name	Rows	Bytes
0	SELECT STATEMENT		1	45
1	TABLE ACCESS BY INDEX ROWID	PERF_TEAM	1	45
2	BITMAP CONVERSION TO ROWIDS			
3	BITMAP AND			
4	BITMAP INDEX SINGLE VALUE	IX_B1		
5	BITMAP INDEX SINGLE VALUE	IX_B2		


Predicate: 4 - access("GENDER"='M') 5 - access("COUNTRY"='FR')



# Bitmap Operations

- BITMAP CONVERSION:
  - TO ROWIDS
  - FROM ROWIDS
  - COUNT
- BITMAP INDEX:
  - SINGLE VALUE
  - RANGE SCAN
  - FULL SCAN
- BITMAP MERGE
- BITMAP AND/OR
- BITMAP MINUS
- BITMAP KEY ITERATION





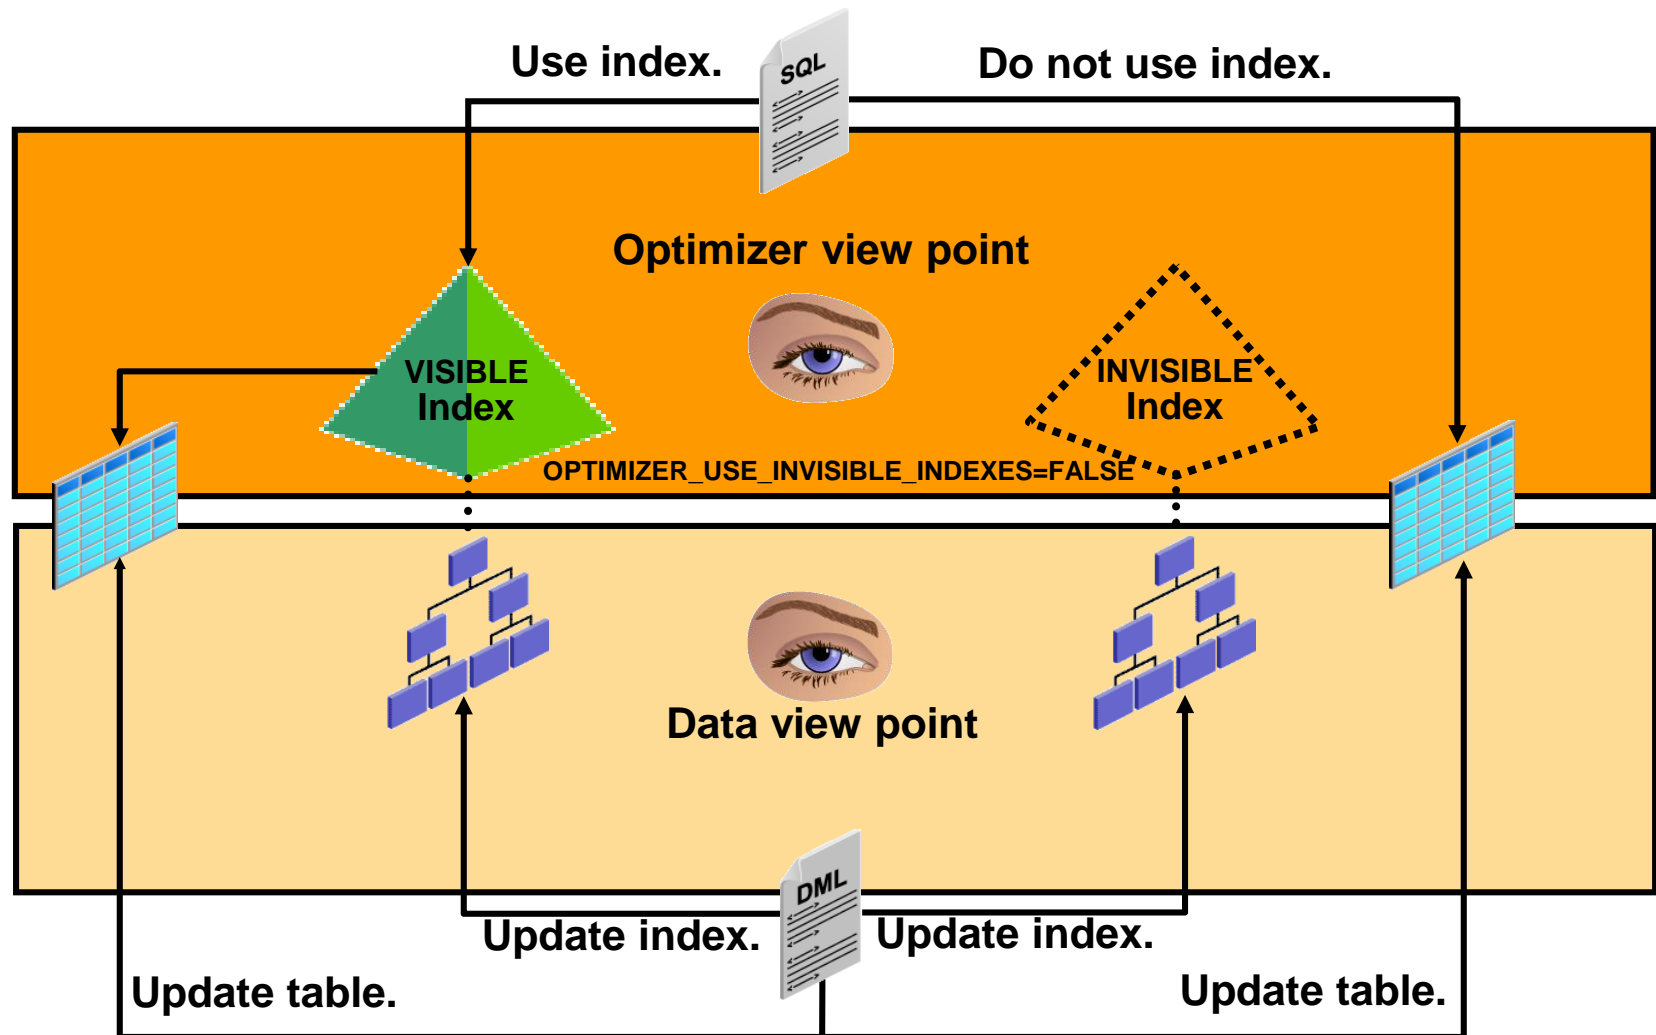
```
create index cars_make_model_idx on
cars (make, model);
```

```
select *
from cars
where make = 'CITROËN' and model = '2CV';
```

Id	Operation	Name
0	SELECT STATEMENT	
1	TABLE ACCESS BY INDEX ROWID	CUSTOMERS
* 2	INDEX RANGE SCAN	CARS_MAKE_MODEL_IDX



## Invisible Index: Overview



## Join Methods



### ➤ A join:

- Defines the relationship between two row sources
- Is a method of combining data from two data sources
- Is controlled by join predicates, which define how the objects are related
- Join methods:
  - Nested loops
  - Sort-merge join
  - Hash join

```
SELECT e.ename, d.dname
FROM dept d JOIN emp e USING (deptno)
WHERE e.job = 'ANALYST' OR e.empno = 9999
```

← Join predicate  
← Nonjoin predicate


```
SELECT e.ename, d.dname
FROM emp e, dept d
WHERE e.deptno = d.deptno AND
(e.job = 'ANALYST' OR e.empno = 9999);
```

← Join predicate  
← Nonjoin predicate



## Nested Loops Join

- Driving row source is scanned.
- Each row returned drives a lookup in inner row source.
- Joining rows are then returned.




```
select ename, e.deptno, d.deptno, d.dname
      from emp e, dept d
     where e.deptno = d.deptno and ename like 'A%';
```

Id	Operation	Name	Rows	Cost
0	SELECT STATEMENT		2	4
1	NESTED LOOPS		2	4
2	TABLE ACCESS FULL	EMP	2	2
3	TABLE ACCESS BY INDEX ROWID	DEPT	1	1
4	INDEX UNIQUE SCAN	PK_DEPT	1	

```
2 - filter("E"."ENAME" LIKE 'A%')
4 - access("E"."DEPTNO"="D"."DEPTNO")
```

## Nested Loops Join: Prefetching



```
select ename, e.deptno, d.deptno, d.dname
      from emp e, dept d
where e.deptno = d.deptno and ename like 'A%';
```


	0		SELECT STATEMENT				2		84		5
	1		TABLE ACCESS BY INDEX ROWID		DEPT		1		22		1
	2		NESTED LOOPS				2		84		5
	* 3		TABLE ACCESS FULL		EMP		2		40		3
	* 4		INDEX RANGE SCAN		IDEPT		1				0

```
3 - filter("E"."ENAME" LIKE 'A%')
4 - access("E"."DEPTNO"="D"."DEPTNO")
```

## Nested Loops Join: 11g Implementation



```
select ename, e.deptno, d.deptno, d.dname
      from emp e, dept d
 where e.deptno = d.deptno and ename like 'A%';
```




Explain Plan x				
0.299 seconds				
OPERATION	OBJECT_NAME	OPTIONS	COST	CARDINALITY
SELECT STATEMENT			4	1
NESTED LOOPS				
NESTED LOOPS			4	1
TABLE ACCESS	EMP	FULL	3	1
Filter Predicates ENAME LIKE 'A%'				
INDEX	PK_DEPT	UNIQUE SCAN	0	1
Access Predicates E.DEPTNO=D.DEPTNO				
TABLE ACCESS	DEPT	BY INDEX RO...	1	1

## Sort Merge Join



- First and second row sources are sorted by the same sort key.
- Sorted rows from both tables are merged.

```
select /*+ USE_MERGE(d e) NO_INDEX(d) */  
  ename, e.deptno, d.deptno, dname  
  from emp e, dept d  
 where e.deptno = d.deptno and ename > 'A'
```




OPERATION	OBJECT_NAME	OPTIONS	COST	CARDINALI...
SELECT STATEMENT			8	14
MERGE JOIN			8	14
SORT		JOIN	4	4
TABLE ACCESS	DEPT	FULL	3	4
SORT		JOIN	4	14
Access Predicates				
AND				
E.DEPTNO=D.DEPTNO				
SYS_OP_DESCEND(E.DEPT				
Filter Predicates				
E.DEPTNO=D.DEPTNO				
TABLE ACCESS	EMP	FULL	3	14
Filter Predicates				
ENAME>'A'				

# Hash Join




- The smallest row source is used to build a hash table.
- The second row source is hashed and checked against the hash table.



```
select /*+ USE_HASH(e d) */
  ename, e.deptno, d.deptno, dname
  from emp e, dept d
 where e.deptno = d.deptno and ename like 'A%'
```

OPERATION	OBJECT_NAME	OPTIONS	COST	CARDINALI...
SELECT STATEMENT			7	1
HASH JOIN			7	1
Access Predicates				
AND				
E.DEPTNO=D.DEPTNO				
SYS_OP_DESCEND(E.DEPTNO)				
TABLE ACCESS	EMP	FULL	3	1
Filter Predicates				
ENAME LIKE 'A%'				
TABLE ACCESS	DEPT	FULL	3	4

## Cartesian Join



```
select ename, e.deptno, d.deptno, dname
from emp e, dept d where ename like 'A%';
```

Explain Plan x

0.295 seconds

OPERATION	OBJECT_NAME	OPTIONS	COST	CARDINALI...
SELECT STATEMENT			6	4
MERGE JOIN		CARTESIAN	6	4
TABLE ACCESS	EMP	FULL	3	1
Filter Predicates				
ENAME LIKE 'A%'				
BUFFER		SORT	3	4
TABLE ACCESS	DEPT	FULL	3	4





## Join Types

- A join operation combines the output from two row sources and returns one resulting row source.
- Join operation types include the following :
  - Join (Equijoin/Natural – Nonequijoin)
  - Outer join (Full, Left, and Right)
  - Semi join: `EXISTS` subquery
  - Anti join: `NOT IN` subquery
  - Star join (Optimization)

## Equijoins and Nonequijoins



```
SELECT e.ename, e.sal, s.grade
FROM   emp e ,salgrade s
WHERE  e.sal = s.hisal;
```

Equijoin

OPERATION	OBJECT_NAME	OPTIONS
SELECT STATEMENT		
HASH JOIN		
Access Predicates E.SAL=S.HISAL		
TABLE ACCESS	SALGRADE	FULL
TABLE ACCESS	EMP	FULL

Nonequijoin

```
SELECT e.ename, e.sal, s.grade
FROM   emp e ,salgrade s
WHERE  e.sal between s.hisal and s.hisal;
```

OPERATION	OBJECT_NAME	OPTIONS
SELECT STATEMENT		
MERGE JOIN		
SORT		JOIN
TABLE ACCESS	SALGRADE	FULL
FILTER		
Filter Predicates E.SAL<=S.HISAL		
SORT		JOIN
Access Predicates E.SAL>=S.HISAL		
Filter Predicates E.SAL>=S.HISAL		
TABLE ACCESS	EMP	FULL

# Outer Joins




➤ An outer join returns a row even if no match is found.

```
SELECT d.deptno,d.dname,e.empno,e.ename
      FROM emp e, dept d
     WHERE e.deptno(+) = d.deptno;
```

OPERATION	OBJECT_NAME	OPTIONS
SELECT STATEMENT		
MERGE JOIN		OUTER
TABLE ACCESS	DEPT	BY INDEX RO.
INDEX	PK_DEPT	FULL SCAN
SORT		JOIN
Access Predicates		
E.DEPTNO(+) = D.DEPTNO		
Filter Predicates		
E.DEPTNO(+) = D.DEPTNO		
TABLE ACCESS	EMP	FULL

**1**



```
SELECT /*+ USE_HASH(e d) */
d.deptno,d.dname,e.empno,e.ename
      FROM emp e, dept d
     WHERE e.deptno(+) = d.deptno;
```

OPERATION	OBJECT_NAME	OPTIONS
SELECT STATEMENT		
HASH JOIN		OUTER
Access Predicates		
E.DEPTNO(+) = D.DEPTNO		
TABLE ACCESS	DEPT	FULL
TABLE ACCESS	EMP	FULL

**2**


```
SELECT /*+ USE_NL(e d) */
d.deptno,d.dname,e.empno,e.ename
      FROM emp e, dept d
     WHERE e.deptno(+) = d.deptno;
```

OPERATION	OBJECT_NAME	OPTIONS
SELECT STATEMENT		
NESTED LOOPS		OUTER
TABLE ACCESS	DEPT	FULL
TABLE ACCESS	EMP	FULL
Filter Predicates		
E.DEPTNO(+) = D.DEPTNO		

**3**



- Semijoins look only for the first match.



```
SELECT deptno, dname
      FROM dept
 WHERE EXISTS (SELECT 1 FROM emp WHERE emp.deptno=dept.deptno);
```



OPERATION	OBJECT_NAME	OPTIONS
SELECT STATEMENT		
MERGE JOIN		SEMI
TABLE ACCESS	DEPT	BY INDEX RO.
INDEX	PK_DEPT	FULL SCAN
SORT		UNIQUE
Access Predicates		
EMP.DEPTNO=DEPT.DEPTNO		
Filter Predicates		
EMP.DEPTNO=DEPT.DEPTNO		
TABLE ACCESS	EMP	FULL



- Reverse of what would have been returned by a join

**SELECT deptno, dname  
FROM dept  
WHERE deptno not in  
(SELECT deptno FROM emp);**

OPERATION	OBJECT_NAME	OPTIONS
SELECT STATEMENT		
MERGE JOIN		ANTI NA JOIN
SORT		BY INDEX RO..
TABLE ACCESS	DEPT	FULL SCAN
INDEX	PK_DEPT	UNIQUE
SORT		
Access Predicates		DEPTNO=DEPTNO
Filter Predicates		DEPTNO=DEPTNO
TABLE ACCESS	EMP	FULL



**SELECT deptno, dname FROM dept  
WHERE deptno IS NOT NULL AND  
deptno NOT IN  
(SELECT /\*+ HASH\_AJ \*/ deptno FROM  
emp WHERE deptno IS NOT NULL);**

OPERATION	OBJECT_NAME	OPTIONS
SELECT STATEMENT		
HASH JOIN		ANTI
Access Predicates		DEPTNO=DEPTNO
TABLE ACCESS	DEPT	FULL
TABLE ACCESS	EMP	FULL
Filter Predicates		DEPTNO IS NOT NULL



## Invisible Indexes: Examples

- Index is altered as not visible to the optimizer:

```
ALTER INDEX ind1 INVISIBLE;
```

- Optimizer considers this index:

```
SELECT /*+ index(TAB1 IND1) */ COL1 FROM  
TAB1 WHERE ...;
```

```
ALTER INDEX ind1 VISIBLE;
```

```
CREATE INDEX IND1 ON TAB1 (COL1) INVISIBLE;
```



## Guidelines for Managing Indexes

- Create indexes after inserting table data.
- Index the correct tables and columns.
- Order index columns for performance.
- Limit the number of indexes for each table.
- Drop indexes that are no longer required.
- Specify the tablespace for each index.
- Consider parallelizing index creation.
- Consider creating indexes with `NOLOGGING`.
- Consider costs and benefits of coalescing or rebuilding indexes.
- Consider cost before disabling or dropping constraints.



## Investigating Index Usage

➤ An index may not be used for one of many reasons:

- There are functions being applied to the predicate.
- There is a data type mismatch.
- Statistics are old.
- The column can contain null.
- Using the index would actually be slower than not using it.





# Optimizer Statistics

- Describe the database and the objects in the database
- Information used by the query optimizer to estimate:
  - Selectivity of predicates
  - Cost of each execution plan
  - Access method, join order, and join method
  - CPU and input/output (I/O) costs
- Refreshing optimizer statistics whenever they are stale is as important as gathering them:
  - Automatically gathered by the system
  - Manually gathered by the user with `DBMS_STATS`



# Types of Optimizer Statistics

- Table statistics:

- Number of rows
- Number of blocks
- Average row length

- Index Statistics:

- B\*-tree level
- Distinct keys
- Number of leaf blocks
- Clustering factor

- System statistics

- I/O performance and utilization
- CPU performance and utilization

- Column statistics

- Basic: Number of distinct values, number of nulls, average length, min, max
- Histograms (data distribution when the column data is skewed)
- Extended statistics



## Table Statistics (`DBA_TAB_STATISTICS`)

- Table statistics are used to determine:
  - Table access cost
  - Join cardinality
  - Join order
- Some of the table statistics gathered are:
  - Row count (`NUM_ROWS`)
  - Block count (`BLOCKS`) *Exact*
  - Average row length (`AVG_ROW_LEN`)
  - Statistics status (`STALE_STATS`)

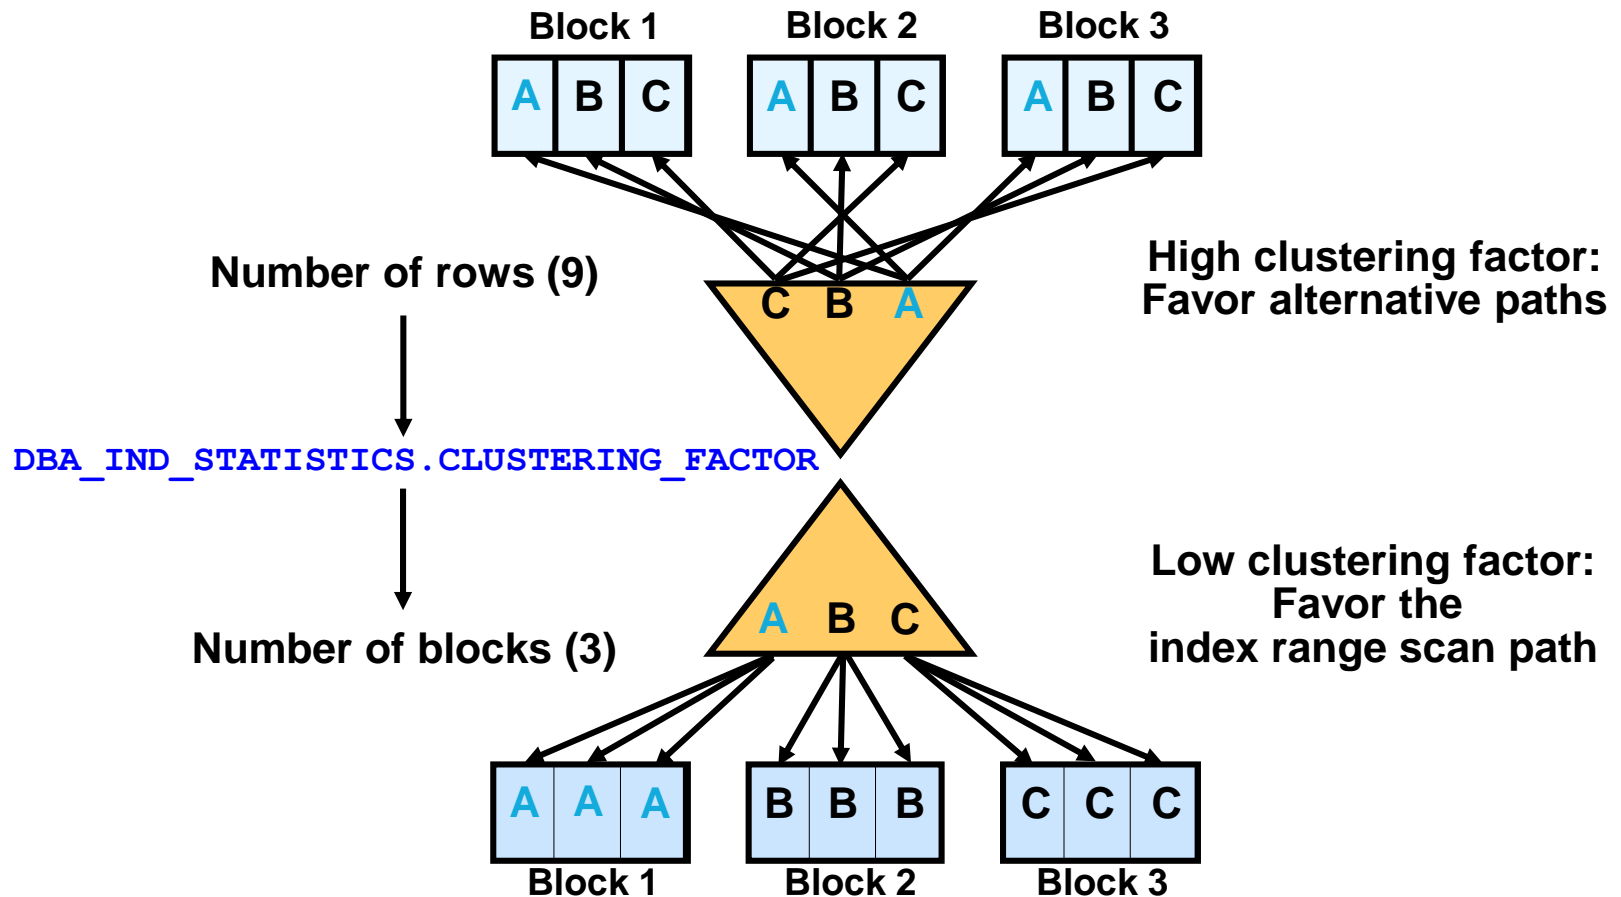


## Index Statistics (DBA\_IND\_STATISTICS)

- Used to decide:
  - Full table scan versus index scan
- Statistics gathered are:
  - B\*-tree level (BLEVEL) *Exact*
  - Leaf block count (LEAF\_BLOCKS)
  - Clustering factor (CLUSTERING\_FACTOR)
  - Distinct keys (DISTINCT\_KEYS)
    - Average number of leaf blocks in which each distinct value in the index appears (AVG\_LEAF\_BLOCKS\_PER\_KEY)
    - Average number of data blocks in the table pointed to by a distinct value in the index (AVG\_DATA\_BLOCKS\_PER\_KEY)
  - Number of rows in the index (NUM\_ROWS)



**Must read all blocks to retrieve all As**



**Only need to read one block to retrieve all As**



## Column Statistics (DBA\_TAB\_COL\_STATISTICS)


- Count of distinct values of the column (NUM\_DISTINCT)
- Low value (LOW\_VALUE) *Exact*
- High value (HIGH\_VALUE) *Exact*
- Number of nulls (NUM\_NULLS)
- Selectivity estimate for nonpopular values (DENSITY)
- Number of histogram buckets (NUM\_BUCKETS)
- Type of histogram (HISTOGRAM)



- The optimizer assumes uniform distributions; this may lead to suboptimal access plans in the case of data skew.
- Histograms:
  - Store additional column distribution information
  - Give better selectivity estimates in the case of nonuniform distributions
- With unlimited resources you could store each different value and the number of rows for that value.
- This becomes unmanageable for a large number of distinct values and a different approach is used:
  - Frequency histogram ( $\text{\#distinct values} \leq \text{\#buckets}$ )
  - Height-balanced histogram ( $\text{\#buckets} < \text{\#distinct values}$ )
- They are stored in `DBA_TAB_HISTOGRAMS`.



10 buckets, 10 distinct values



Distinct values: 1, 3, 5, 7, 10, 16, 27, 32, 39, 49

Number of rows: 40001



## Viewing Frequency Histograms



```
BEGIN
DBMS_STATS.gather_table_STATS (OWNNAME=>'OE', TABNAME=>'INVENTORIES',
METHOD_OPT => 'FOR COLUMNS SIZE 20 warehouse_id');
END;
```

```
SELECT column_name, num_distinct, num_buckets, histogram
FROM    USER_TAB_COL_STATISTICS
WHERE   table_name = 'INVENTORIES' AND
        column_name = 'WAREHOUSE_ID';
```

COLUMN_NAME	NUM_DISTINCT	NUM_BUCKETS	HISTOGRAM
WAREHOUSE_ID	9	9	FREQUENCY


```
SELECT endpoint_number, endpoint_value
FROM    USER_HISTOGRAMS
WHERE   table_name = 'INVENTORIES' and column_name = 'WAREHOUSE_ID'
ORDER BY endpoint_number;
```

ENDPOINT_NUMBER	ENDPOINT_VALUE
36	1
213	2
261	3

...



**5 buckets, 10 distinct values  
(8000 rows per bucket)**



**Distinct values: 1, 3, 5, 7, 10, 16, 27, 32, 39, 49**

**Number of rows: 40001**

## Viewing Height-Balanced Histograms



```
BEGIN
DBMS_STATS.gather_table_STATS(OWNNAME =>'OE', TABNAME=>'INVENTORIES',
METHOD_OPT => 'FOR COLUMNS SIZE 10 quantity_on_hand');
END;
```

```
SELECT column_name, num_distinct, num_buckets, histogram
FROM USER_TAB_COL_STATISTICS
WHERE table_name = 'INVENTORIES' AND column_name = 'QUANTITY_ON_HAND';
```

COLUMN_NAME	NUM_DISTINCT	NUM_BUCKETS	HISTOGRAM
QUANTITY_ON_HAND	237	10	HEIGHT BALANCED


```
SELECT endpoint_number, endpoint_value
FROM USER_HISTOGRAMS
WHERE table_name = 'INVENTORIES' and column_name = 'QUANTITY_ON_HAND'
ORDER BY endpoint_number;
```

ENDPOINT_NUMBER	ENDPOINT_VALUE
0	0
1	27
2	42
3	57
...	



## Histogram Considerations

- Histograms are useful when you have a high degree of skew in the column distribution.
- Histograms are *not* useful for:
  - Columns which do not appear in the `WHERE` or `JOIN` clauses
  - Columns with uniform distributions
  - Equality predicates with unique columns
- The maximum number of buckets is the least (254, # distinct values).
- Do not use histograms unless they substantially improve performance.



$S(\text{MAKE} \wedge \text{MODEL}) = S(\text{MAKE}) \times S(\text{MODEL})$


```
select  
dbms_stats.create_extended_stats('jfv','vehicle','(make,model)')  
from dual;
```

2

```
exec dbms_stats.gather_table_stats('jfv','vehicle',-  
method_opt=>'for all columns size 1 for columns (make,model) size 3')
```

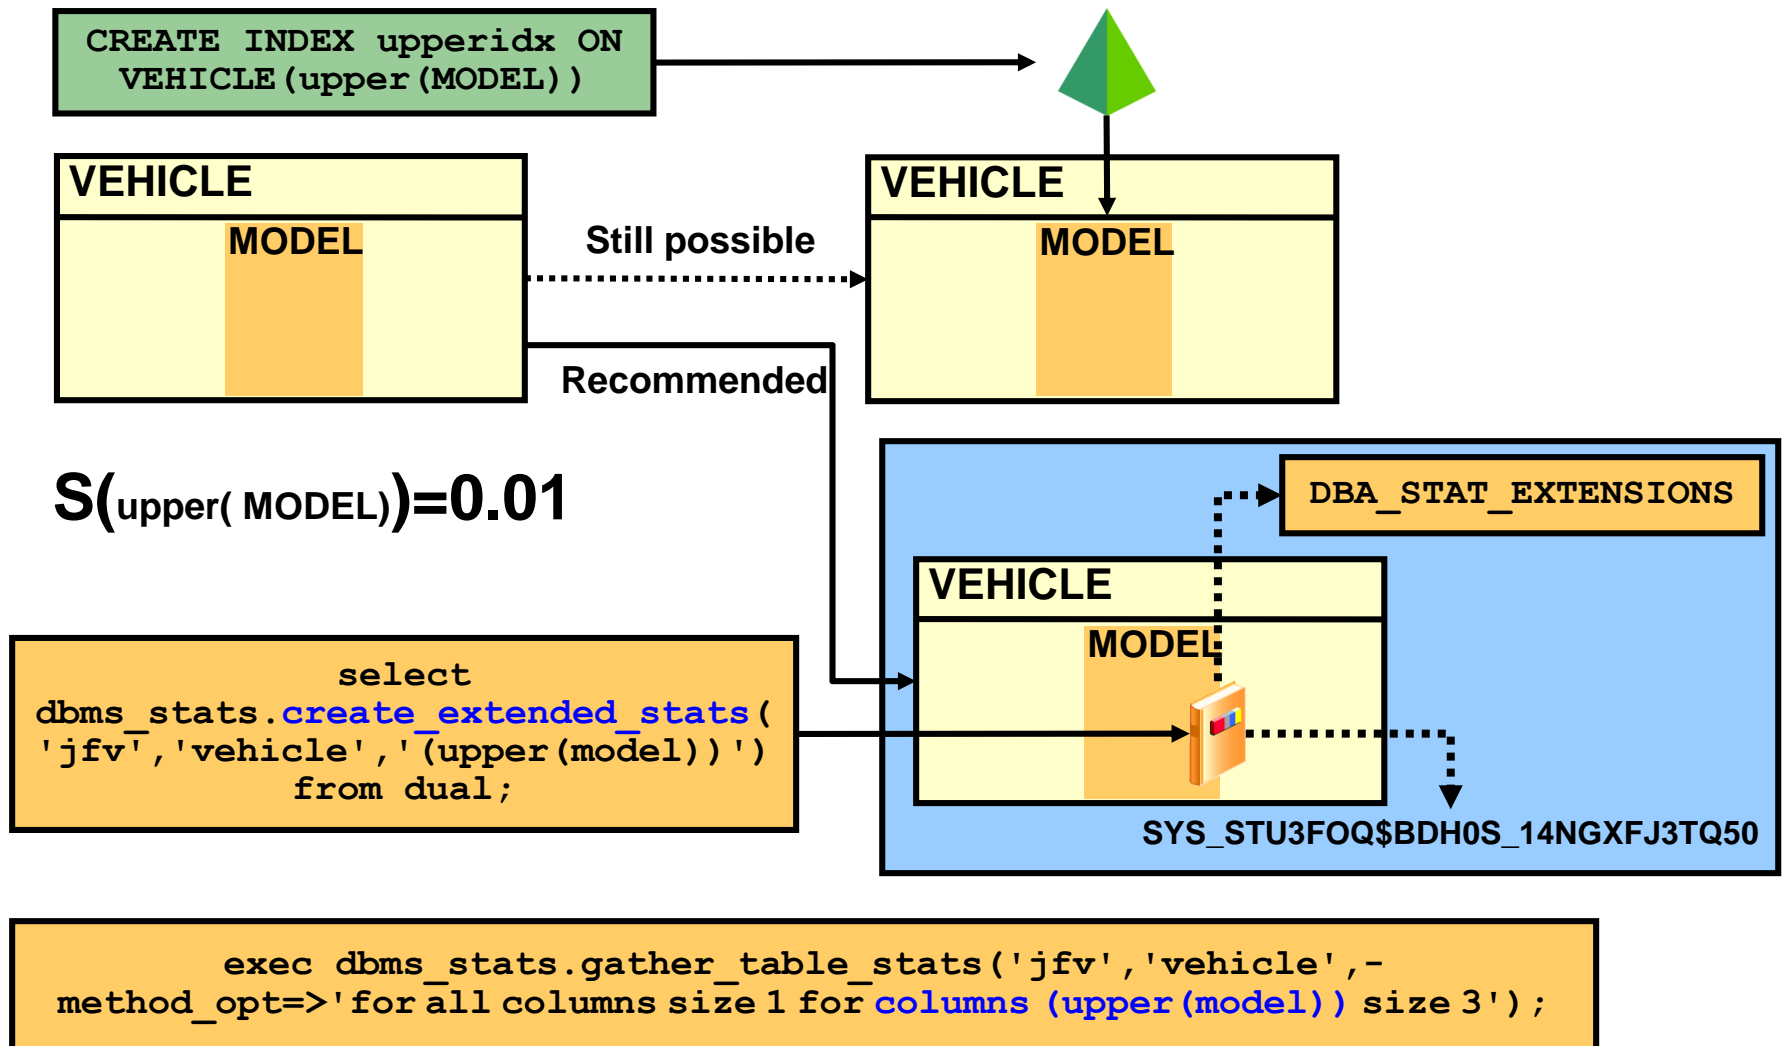
3

DBA\_STAT\_EXTENSIONS



$S(\text{MAKE} \wedge \text{MODEL}) = S(\text{MAKE}, \text{MODEL})$

SYS\_STUF3GLKIOP5F4B0BTTCFTMX0W





# Gathering System Statistics

- System statistics enable the CBO to use CPU and I/O characteristics.
- System statistics must be gathered on a regular basis; this does not invalidate cached plans.
- Gathering system statistics equals analyzing system activity for a specified period of time:
- Procedures:
  - `DBMS_STATS.GATHER_SYSTEM_STATS`
  - `DBMS_STATS.SET_SYSTEM_STATS`
  - `DBMS_STATS.GET_SYSTEM_STATS`
- `GATHERING_MODE`:
  - `NOWORKLOAD|INTERVAL`
  - `START|STOP`



## When to Gather Statistics Manually

- Rely mostly on automatic statistics collection:
  - Change the frequency of automatic statistics collection to meet your needs.
  - Remember that `STATISTICS_LEVEL` should be set to `TYPICAL` or `ALL` for automatic statistics collection to work properly.
- Gather statistics manually for:
  - Objects that are volatile
  - Objects modified in batch operations: Gather statistics as part of the batch operation.
  - External tables, system statistics, fixed objects
  - New objects: Gather statistics right after object creation.





# Mechanisms for Gathering Statistics

- Automatic statistics gathering
  - `gather_stats_prog` automated task
- Manual statistics gathering
  - `DBMS_STATS` package
- Dynamic sampling
- When statistics are missing:

Selectivity:	
Equality	1%
Inequality	5%
Other predicates	5%
Table row length	20
# of index leaf blocks	25
# of distinct values	100
Table cardinality	100
Remote table cardinality	2000



## Manual Statistics Collection: Factors

- Monitor objects for DMLs.
- Determine the correct sample sizes.
- Determine the degree of parallelism.
- Determine if histograms should be used.
- Determine the cascading effects on indexes.
- Procedures to use in `DBMS_STATS`:
  - `GATHER_INDEX_STATS`
  - `GATHER_TABLE_STATS`
  - `GATHER_SCHEMA_STATS`
  - `GATHER_DICTIONARY_STATS`
  - `GATHER_DATABASE_STATS`
  - `GATHER_SYSTEM_STATS`



```
        dbms_stats.gather_table_stats
        ('sh'                -- schema
        , 'customers'        -- table
        , null                -- partition
        , 20                  -- sample size(%)
        , false               -- block sample?
        , 'for all columns'   -- column spec
, 4                          -- degree of parallelism
        , 'default'          -- granularity
        , true );            -- cascade to indexes
```

```
        dbms_stats.set_param('CASCADE',
                              'DBMS_STATS.AUTO CASCADE');
dbms_stats.set_param('ESTIMATE_PERCENT', '5');
        dbms_stats.set_param('DEGREE', 'NULL');
```



- Prevents automatic gathering
- Is mainly used for volatile tables:
  - Lock without statistics implies dynamic sampling.
- Lock with statistics for representative values.

- The `FORCE`

```
BEGIN
DBMS_STATS.DELETE_TABLE_STATS('OE','ORDERS');
DBMS_STATS.LOCK_TABLE_STATS('OE','ORDERS');
END;
```

```
BEGIN
DBMS_STATS.GATHER_TABLE_STATS('OE','ORDERS');
DBMS_STATS.LOCK_TABLE_STATS('OE','ORDERS');
END;
```

```
SELECT stattype_locked FROM
       dba_tab_statistics;
```



- Past Statistics may be restored with `DBMS_STATS.RESTORE_*_STATS` procedures

- Sta  
· )  
· )  
▪ Sta  
· )  
· )


```
BEGIN
    DBMS_STATS.RESTORE_TABLE_STATS(
        OWNNAME=>'OE', TABNAME=>'INVENTORIES',
        AS_OF_TIMESTAMP=>'15-JUL-10 09.28.01.597526000 AM -05:00');
END;
```



## Export and Import Statistics

### ➤ Use DBMS\_STATS procedures:

- `CREATE_STAT_TABLE` creates the statistics table.
- `EXPORT_*_STATS` moves the statistics to the statistics table.
- Use Data Pump to move the statistics table.
- `IMPORT_*_STATS` moves the statistics to data dictionary.





## Optimizer Hints: Overview

### ➤ Optimizer hints:

- Influence optimizer decisions
- Example:

- HINTS SHOULD ONLY BE USED AS A LAST RESORT.
- When you use a hint, it is good practice to also add a comment about that hint

```
SELECT /*+ INDEX(e empfirstname_idx) skewed col */ *  
      FROM employees e  
      WHERE first_name='David'
```



<b>Single-table hints</b>	Specified on one table or view
<b>Multitable hints</b>	Specify more than one table or view
<b>Query block hints</b>	Operate on a single query block
<b>Statement hints</b>	Apply to the entire SQL statement

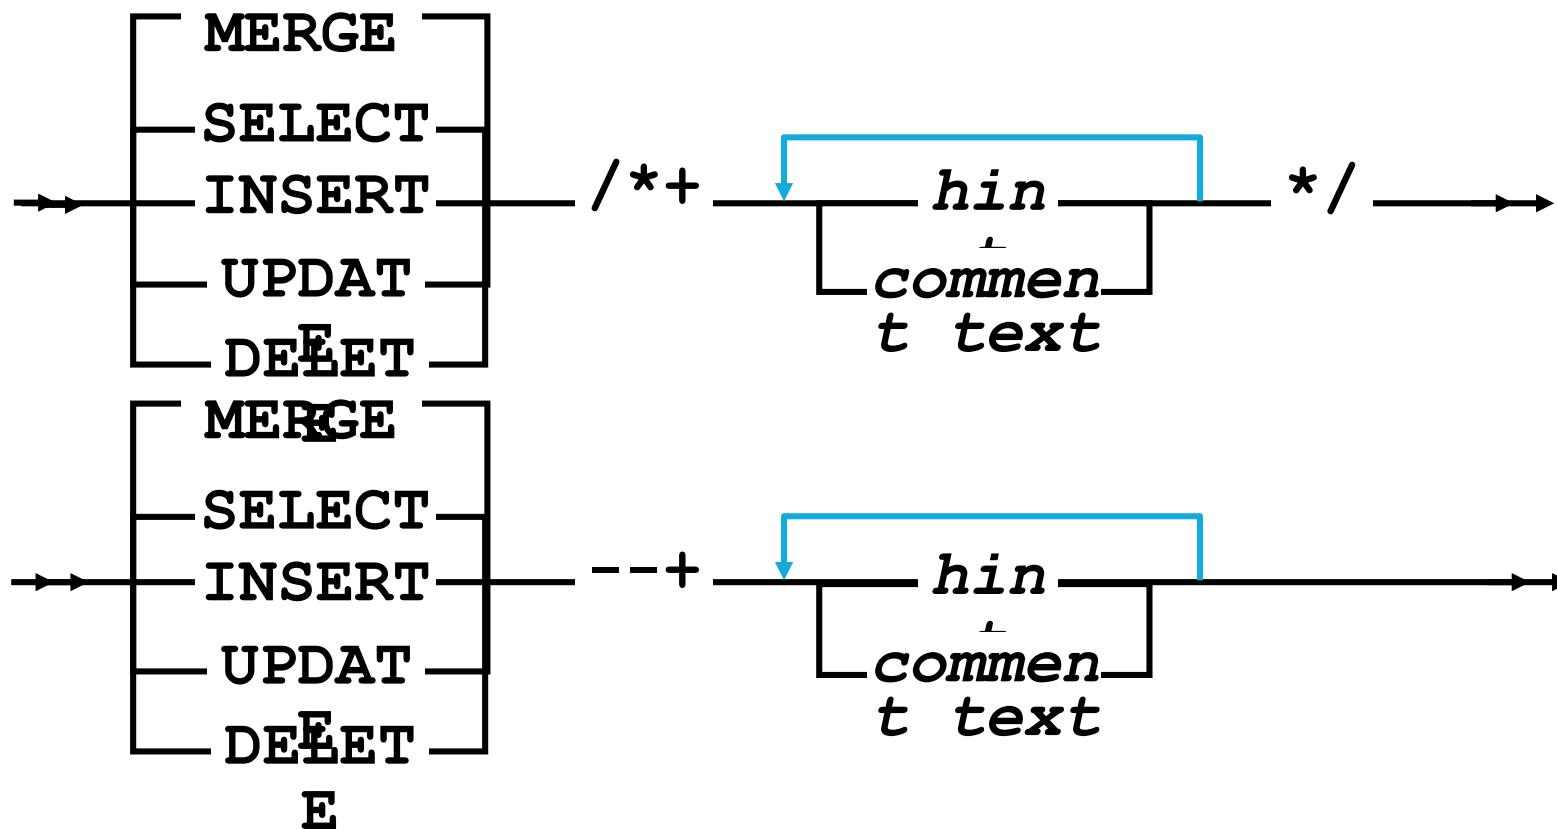


## Specifying Hints



➤ Hints apply to the optimization of only one statement block:

- A self-contained DML statement against a table
- A top-level DML or a subquery





## Rules for Hints

- Place hints immediately after the first SQL keyword of a statement block.
- Each statement block can have only one hint comment, but it can contain multiple hints.
- Hints apply to only the statement block in which they appear.
- If a statement uses aliases, hints must reference the aliases rather than the table names.
- The optimizer ignores hints specified incorrectly without raising errors.



## Hint Recommendations

- Use hints carefully because they imply a high-maintenance load.
- Be aware of the performance impact of hard-coded hints when they become less valid.



```
UPDATE /*+ INDEX(p PRODUCTS_PROD_CAT_IX) */  
      products p  
      SET   p.prod_min_price =  
            (SELECT  
              (pr.prod_list_price*.95)  
              FROM products pr  
              WHERE p.prod_id = pr.prod_id)  
      WHERE p.prod_category = 'Men'  
AND    p.prod_status = 'available, on stock'  
      /
```



## Hint Categories

- There are hints for:
  - Optimization approaches and goals
  - Access paths
  - Query transformations
  - Join orders
  - Join operation
  - Parallel execution
  - Additional hints



<b>ALL_ROWS</b>	Selects a cost-based approach with a goal of best throughput
<b>FIRST_ROWS (<i>n</i>)</b>	Instructs the Oracle server to optimize an individual SQL statement for fast response

➤ **Note:** The `ALTER SESSION... SET OPTIMIZER_MODE` statement does not affect SQL that is run from within PL/SQL.



<b>FULL</b>	Performs a full table scan
<b>CLUSTER</b>	Accesses table using a cluster scan
<b>HASH</b>	Accesses table using a hash scan
<b>ROWID</b>	Accesses a table by ROWID
<b>INDEX</b>	Selects an index scan for the specified table
<b>INDEX_ASC</b>	Scans an index in ascending order
<b>INDEX_COMBINE</b>	Explicitly chooses a bitmap access path



<b>INDEX_JOIN</b>	Instructs the optimizer to use an index join as an access path
<b>INDEX_DESC</b>	Scans an index in descending order
<b>INDEX_FFS</b>	Performs a fast-full index scan
<b>INDEX_SS</b>	Performs an index skip scan
<b>NO_INDEX</b>	Does not allow using a set of indexes





```
SELECT /*+INDEX_COMBINE(CUSTOMERS)*/  
        cust_last_name  
        FROM SH.CUSTOMERS  
        WHERE ( CUST_GENDER= 'F' AND  
                CUST_MARITAL_STATUS = 'single')  
OR       CUST_YEAR_OF_BIRTH BETWEEN '1917'  
        AND '1920';
```

## The INDEX\_COMBINE Hint: Example



### Execution Plan

```
-----  
      |      0 | SELECT STATEMENT                                | |
      |      1 | TABLE ACCESS BY INDEX ROWID | CUSTOMERS      |  
      |      2 | BITMAP CONVERSION TO ROWIDS  |                |  
      |      3 | BITMAP OR                    |                |  
      |      4 | BITMAP MERGE                 |                |  
      |      5 | BITMAP INDEX RANGE SCAN      | CUST_YOB_BIX   |  
      |      6 | BITMAP AND                   |                |  
      |      7 | BITMAP INDEX SINGLE VALUE    | CUST_MARITAL_BIX |  
      |      8 | BITMAP INDEX SINGLE VALUE    | CUST_GENDER_BIX |
```



<b>NO_QUERY_TRANSFORMATION</b>	Skips all query transformation
<b>USE_CONCAT</b>	Rewrites OR into UNION ALL and disables INLIST processing
<b>NO_EXPAND</b>	Prevents OR expansions
<b>REWRITE</b>	Rewrites query in terms of materialized views
<b>NO_REWRITE</b>	Turns off query rewrite
<b>UNNEST</b>	Merges subquery bodies into surrounding query block
<b>NO_UNNEST</b>	Turns off unnesting



<b>MERGE</b>	Merges complex views or subqueries with the surrounding query
<b>NO_MERGE</b>	Prevents merging of mergeable views
<b>STAR_TRANSFORMATION</b>	Makes the optimizer use the best plan in which the transformation can be used
<b>FACT</b>	Indicates that the hinted table should be considered as a fact table
<b>NO_FACT</b>	Indicates that the hinted table should not be considered as a fact table



<b>ORDERED</b>	Causes the Oracle server to join tables in the order in which they appear in the FROM clause
<b>LEADING</b>	Uses the specified tables as the first table in the join order



<b>USE_NL</b>	Joins the specified table using a nested loop join
<b>NO_USE_NL</b>	Does not use nested loops to perform the join
<b>USE_NL_WITH_INDEX</b>	Similar to <code>USE_NL</code> , but must be able to use an index for the join
<b>USE_MERGE</b>	Joins the specified table using a sort-merge join
<b>NO_USE_MERGE</b>	Does not perform sort-merge operations for the join
<b>USE_HASH</b>	Joins the specified table using a hash join
<b>NO_USE_HASH</b>	Does not use hash join
<b>DRIVING_SITE</b>	Instructs the optimizer to execute the query at a different site than that selected by the database



<b>APPEND</b>	Enables direct-path <code>INSERT</code>
<b>NOAPPEND</b>	Enables regular <code>INSERT</code>
<b>CURSOR_SHARING_EXACT</b>	Prevents replacing literals with bind variables
<b>CACHE</b>	Overrides the default caching specification of the table
<b>PUSH_PRED</b>	Pushes join predicate into view
<b>PUSH_SUBQ</b>	Evaluates nonmerged subqueries first
<b>DYNAMIC_SAMPLING</b>	Controls dynamic sampling to improve server performance



<b>MONITOR</b>	Forces real-time query monitoring
<b>NO_MONITOR</b>	Disables real-time query monitoring
<b>RESULT_CACHE</b>	Caches the result of the query or query fragment
<b>NO_RESULT_CACHE</b>	Disables result caching for the query or query fragment
<b>OPT_PARAM</b>	Sets initialization parameter for query duration





## Hints and Views

- Do not use hints in views.
- Use view-optimization techniques:
  - Statement transformation
  - Results accessed like a table
- Hints can be used on mergeable views and nonmergeable views.



- Extended hint syntax enables specifying for tables that appear in views
- References a table name in the hint with a recursive dot notation

```
CREATE view city_view AS
  SELECT *
    FROM customers c
 WHERE cust_city like 'S%';
```

```
SELECT /*+ index(v.c cust_credit_limit_idx) */
       v.cust_last_name, v.cust_credit_limit
    FROM city_view v
 WHERE cust_credit_limit > 5000;
```



## Specifying a Query Block in a Hint

```
explain plan for
select /*+ FULL(@strange dept) */ ename
from emp e, (select /*+ QB_NAME(strange) */ *
             from dept where deptno=10) d
where e.deptno = d.deptno and d.loc = 'C';
```

```
SELECT * FROM TABLE(DBMS_XPLAN.DISPLAY(NULL, NULL, 'ALL'));
```

Plan hash value: 615168685

-----						
Id		Operation	Name	Rows	Bytes	Cost(%CPU)
-----						
0		SELECT STATEMENT		1	41	7 (15)
* 1		HASH JOIN		1	41	7 (15)
* 2		TABLE ACCESS FULL	DEPT	1	21	3 (0)
* 3		TABLE ACCESS FULL	EMP	3	60	3 (0)
-----						

Query Block Name / Object Alias (identified by operation id):

-----  
1 - SEL\$DB579D14  
2 - SEL\$DB579D14 / DEPT@STRANGE  
3 - SEL\$DB579D14 / E@SEL\$1



```
SELECT /*+ LEADING(e2 e1) USE NL(e1)
         INDEX(e1 emp_emp_id_pk) USE MERGE(j) FULL(j) */
       e1.first_name, e1.last_name, j.job_id,
       sum(e2.salary) total_sal
FROM   hr.employees e1, hr.employees e2,
       hr.job_history j
WHERE  e1.employee_id = e2.manager_id
AND    e1.employee_id = j.employee_id
AND    e1.hire_date = j.start_date
GROUP BY e1.first_name, e1.last_name, j.job_id
ORDER BY total_sal;
```



# Materialized Views



- After completing this lesson, you should be able to do the following:
  - Identify the characteristics and benefits of materialized views
  - Use materialized views to enable query rewrites
  - Verify the properties of materialized views
  - Perform refreshes on materialized views



➤ A materialized view:

- Is a precomputed set of results
- Has its own data segment and offers:
  - Space management options
  - Use of its own indexes
- Is useful for:
  - Expensive and complex joins
  - Summary and aggregate data



```
SELECT c.cust_id, SUM(amount_sold)
FROM   sales s, customers c
WHERE  s.cust_id = c.cust_id
GROUP BY c.cust_id;
```



```
CREATE TABLE cust_sales_sum AS
SELECT c.cust_id, SUM(amount_sold) AS amount
FROM   sales s, customers c
WHERE  s.cust_id = c.cust_id
GROUP BY c.cust_id;
```



```
SELECT * FROM cust_sales_sum;
```





```
CREATE MATERIALIZED VIEW cust_sales_mv
ENABLE QUERY REWRITE AS
SELECT c.cust_id, SUM(amount_sold) AS amount
FROM   sales s, customers c
WHERE  s.cust_id = c.cust_id
GROUP BY c.cust_id;
```

↓

```
SELECT c.cust_id, SUM(amount_sold)
FROM   sales s, customers c
WHERE  s.cust_id = c.cust_id
GROUP BY c.cust_id;
```

↓

### Execution Plan

```
-----
0      SELECT STATEMENT Optimizer=ALL_ROWS (Cost=6 ...)

1      0      MAT VIEW REWRITE ACCESS (FULL) OF 'CUST_SALES_MV' (MAT_VIEW
              REWRITE) (Cost=6 ...)
```



## How Many Materialized Views?

- One materialized view for multiple queries:
  - One materialized view can be used to satisfy multiple queries.
  - Less disk space is needed.
  - Less time is needed for maintenance.
- Query rewrite chooses the materialized view to use.
- One materialized view per query:
  - Is not recommended because it consumes too much disk space
  - Improves one query's performance

## Creating Materialized Views: Syntax Options



```
CREATE MATERIALIZED VIEW mview_name
  [TABLESPACE ts_name]
  [PARALLEL (DEGREE n)]
  [BUILD {IMMEDIATE|DEFERRED}]
  [{ REFRESH {FAST|COMPLETE|FORCE}
    | {ON COMMIT|ON DEMAND}}
   | NEVER REFRESH ]
  [{ENABLE|DISABLE} QUERY REWRITE]

AS SELECT ... FROM ...
```



```
CREATE MATERIALIZED VIEW cost_per_year_mv
ENABLE QUERY REWRITE
AS
SELECT      t.week_ending_day
,           t.calendar_year
,           p.prod_subcategory
,           sum(c.unit_cost) AS dollars
FROM        costs c
,           times t
,           products p
WHERE       c.time_id = t.time_id
AND         c.prod_id = p.prod_id
GROUP BY    t.week_ending_day
,           t.calendar_year
,           p.prod_subcategory;

Materialized view created.
```



- Materialized views with aggregates

```
CREATE MATERIALIZED VIEW cust_sales_mv AS
SELECT c.cust_id, s.channel_id,
       SUM(amount_sold) AS amount
FROM   sales s, customers c
WHERE  s.cust_id = c.cust_id
GROUP BY c.cust_id, s.channel_id;
```

```
CREATE MATERIALIZED VIEW sales_products_mv AS
SELECT s.time_id, p.prod_name
FROM   sales s, products p
WHERE  s.prod_id = p.prod_id;
```



- You can specify how you want your materialized views to be refreshed from the detail tables by selecting one of four options:
  - `COMPLETE`
  - `FAST`
  - `FORCE`
  - `NEVER`
- You can view the `REFRESH_METHOD` in the `ALL_MVIEWS` data dictionary view.



- Manual refresh
  - Specify `ON DEMAND` option
  - By using the `DBMS_MVIEW` package
- Automatic refresh Synchronous
  - Specify `ON COMMIT` option
  - Upon commit of changes to the underlying tables but independent of the committing transaction
- Automatic refresh Asynchronous
  - Specify using `START WITH` and `NEXT` clauses
  - Defines a refresh interval for the materialized view
- `REFRESH_MODE` in `ALL_MVIEWS`



## Manual Refresh with `DBMS_MVIEW`

- For `ON DEMAND` refresh only
- Three procedures with the `DBMS_MVIEW` package:
  - `REFRESH`
  - `REFRESH_ALL_MVIEWS`
  - `REFRESH_DEPENDENT`





Specific materialized views:

```
Exec DBMS_MVIEW.REFRESH('cust_sales_mv');
```

### Materialized views based on one or more tables:

```
VARIABLE fail NUMBER;  
exec DBMS_MVIEW.REFRESH_DEPENDENT(-  
:fail, 'CUSTOMERS, SALES');
```

### All materialized views due for refresh:

```
VARIABLE fail NUMBER;  
exec DBMS_MVIEW.REFRESH_ALL_MVIEWS(:fail);
```



- If you want to use a materialized view instead of the base tables, a query must be rewritten.
- Query rewrites are transparent to applications.
- Query rewrites do not require special privileges on the materialized view.
- A materialized view can be enabled or disabled for query rewrites.



- Use `EXPLAIN PLAN` or `AUTOTRACE` to verify that query rewrites occur.
- Check the query response:
  - Fewer blocks are accessed.
  - Response time should be significantly better.

## Enabling and Controlling Query Rewrites



- Query rewrites are available with cost-based optimization only.

- The following parameters control query rewrites:

```
QUERY_REWRITE_ENABLED = {true|false|force}  
QUERY_REWRITE_INTEGRITY =  
{enforced|trusted|stale_tolerated}
```



```
EXPLAIN PLAN FOR
SELECT    t.week_ending_day
,         t.calendar_year
,         p.prod_subcategory
,         sum(c.unit_cost) AS dollars
FROM      costs c
,         times t
,         products p
WHERE     c.time_id = t.time_id
. . .
```

### Execution Plan


```
-----
0      SELECT STATEMENT Optimizer=ALL_ROWS (Cost...)
1      0      MAT_VIEW REWRITE ACCESS (FULL) OF 'costs_per_year_mv' (
              MAT_VIEW REWRITE) (Cost...)
```

## Query Rewrite: Example



```
SELECT    t.week_ending_day
,         t.calendar_year
,         p.prod_subcategory
,         sum(c.unit_cost) AS dollars
FROM      costs c, times t, products p
WHERE     c.time_id = t.time_id
AND       c.prod_id = p.prod_id
AND       t.calendar_year = '1999'
GROUP BY  t.week_ending_day, t.calendar_year
,         p.prod_subcategory
HAVING    sum(c.unit_cost) > 10000;
```

```
SELECT    week_ending_day
,         prod_subcategory
,         dollars
FROM      cost_per_year_mv
WHERE     calendar_year = '1999'
AND       dollars > 10000;
```

A blue arrow pointing downwards from the first query block to the second query block.



```
CREATE MATERIALIZED VIEW cust_orders_mv
ENABLE QUERY REWRITE AS
SELECT c.customer_id, SUM(order_total) AS amt
FROM   oe.orders s, oe.customers c
WHERE  s.customer_id = c.customer_id
GROUP BY c.customer_id;
```



```
SELECT /*+ REWRITE_OR_ERROR */ c.customer_id,
SUM(order_total) AS amt
FROM   oe.orders s, oe.customers c
WHERE  s.customer_id = c.customer_id
GROUP BY c.customer_id;
```



```
ORA-30393: a query block in the statement did
not rewrite
```



- For a given workload, the SQL Access Advisor:
  - Recommends creating the appropriate:
    - Materialized views
    - Materialized view logs
    - Indexes
  - Provides recommendations to optimize for:
    - Fast refresh
    - Query rewrite
  - Can be run:
    - From Oracle Enterprise Manager by using the SQL Access Advisor Wizard
    - By invoking the `DBMS_ADVISOR` package



# Using the DBMS\_MVIEW Package



## ➤ DBMS\_MVIEW methods

- EXPLAIN\_MVIEW
- EXPLAIN\_REWRITE
- TUNE\_MVIEW

## Tuning Materialized Views for Fast Refresh and Query Rewrite



```
DBMS_ADVISOR.TUNE_MVIEW (  
    task_name IN OUT VARCHAR2,  
    mv_create_stmt IN [CLOB | VARCHAR2]  
);
```

## Results of Tune\_MVIEW



- IMPLEMENTATION recommendations
  - CREATE MATERIALIZED VIEW LOG statements
  - ALTER MATERIALIZED VIEW LOG FORCE statements
  - One or more CREATE MATERIALIZED VIEW statements
- UNDO recommendations
  - DROP MATERIALIZED VIEW statements



## DBMS\_MVIEW.EXPLAIN\_MVIEW Procedure

- Accepts:
  - Materialized view name
  - SQL statement
- Advises what is and what is not possible:
  - For an existing materialized view
  - For a potential materialized view before you create it
- Stores results in `MV_CAPABILITIES_TABLE` (relational table) or in a `VARRAY`
- `utlxmlv.sql` must be executed as the current user to create `MV_CAPABILITIES_TABLE`.

## Explain Materialized View: Example



```
EXEC dbms_mview.explain_mview (  
    'cust_sales_mv', '123');
```

```
SELECT capability_name, possible, related_text,msgtxt  
FROM mv_capabilities_table  
WHERE statement_id = '123' ORDER BY seq;
```

CAPABILITY_NAME	P	RELATED_TE	MSGTXT
-----	-	-----	-----
...			
REFRESH_COMPLETE	Y		
REFRESH_FAST	N		
REWRITE	N		
PCT_TABLE	N	SALES	no partition key or PMARKER in select list
PCT_TABLE	N	CUSTOMERS	relation is not a partitioned table
...			



## ➤ Query rewrite considerations:

- Constraints
- Outer joins
- Text match
- Aggregates
- Grouping conditions
- Expression matching
- Date folding
- Statistics



<b>REWRITE</b>	<b>Rewrites a query in terms of materialized views</b>
<b>REWRITE_OR_ERROR</b>	<b>Forces an error if a query rewrite is not possible</b>
<b>NO_REWRITE</b>	<b>Disables query rewrite for the query block</b>

# Summary



- In this lesson, you should have learned how to:
  - Create materialized views
  - Enable query rewrites using materialized views





- After completing this lesson, you should be able to do the following:
  - Identify the characteristics and benefits of materialized views
  - Use materialized views to enable query rewrites
  - Verify the properties of materialized views
  - Perform refreshes on materialized views



# Materialized Views

- A materialized view:
  - Is a precomputed set of results
  - Has its own data segment and offers:
    - Space management options
    - Use of its own indexes
  - Is useful for:
    - Expensive and complex joins
    - Summary and aggregate data



```
SELECT c.cust_id, SUM(amount_sold)
FROM   sales s, customers c
WHERE  s.cust_id = c.cust_id
GROUP BY c.cust_id;
```



```
CREATE TABLE cust_sales_sum AS
SELECT c.cust_id, SUM(amount_sold) AS amount
FROM   sales s, customers c
WHERE  s.cust_id = c.cust_id
GROUP BY c.cust_id;
```



```
SELECT * FROM cust_sales_sum;
```



```
CREATE MATERIALIZED VIEW cust_sales_mv
ENABLE QUERY REWRITE AS
SELECT c.cust_id, SUM(amount_sold) AS amount
FROM   sales s, customers c
WHERE  s.cust_id = c.cust_id
GROUP BY c.cust_id;
```

↓

```
SELECT c.cust_id, SUM(amount_sold)
FROM   sales s, customers c
WHERE  s.cust_id = c.cust_id
GROUP BY c.cust_id;
```

↓

### Execution Plan

```
-----
0      SELECT STATEMENT Optimizer=ALL_ROWS (Cost=6 ...)

1      0      MAT VIEW REWRITE ACCESS (FULL) OF 'CUST_SALES_MV' (MAT_VIEW
              REWRITE) (Cost=6 ...)
```



## How Many Materialized Views?

- One materialized view for multiple queries:
  - One materialized view can be used to satisfy multiple queries.
  - Less disk space is needed.
  - Less time is needed for maintenance.
- Query rewrite chooses the materialized view to use.
- One materialized view per query:
  - Is not recommended because it consumes too much disk space
  - Improves one query's performance

## Creating Materialized Views: Syntax Options



```
CREATE MATERIALIZED VIEW mview_name
  [TABLESPACE ts_name]
  [PARALLEL (DEGREE n)]
  [BUILD {IMMEDIATE|DEFERRED}]
  [{ REFRESH {FAST|COMPLETE|FORCE}
    | {ON COMMIT|ON DEMAND}}
   | NEVER REFRESH ]
  [{ENABLE|DISABLE} QUERY REWRITE]

AS SELECT ... FROM ...
```



```
CREATE MATERIALIZED VIEW cost_per_year_mv
ENABLE QUERY REWRITE
AS
SELECT    t.week_ending_day
,         t.calendar_year
,         p.prod_subcategory
,         sum(c.unit_cost) AS dollars
FROM      costs c
,         times t
,         products p
WHERE     c.time_id = t.time_id
AND       c.prod_id = p.prod_id
GROUP BY  t.week_ending_day
,         t.calendar_year
,         p.prod_subcategory;

Materialized view created.
```



- Materialized views with aggregates

```
CREATE MATERIALIZED VIEW cust_sales_mv AS
SELECT c.cust_id, s.channel_id,
       SUM(amount_sold) AS amount
FROM   sales s, customers c
WHERE  s.cust_id = c.cust_id
GROUP BY c.cust_id, s.channel_id;
```

```
CREATE MATERIALIZED VIEW sales_products_mv AS
SELECT s.time_id, p.prod_name
FROM   sales s, products p
WHERE  s.prod_id = p.prod_id;
```





- You can specify how you want your materialized views to be refreshed from the detail tables by selecting one of four options:
  - `COMPLETE`
  - `FAST`
  - `FORCE`
  - `NEVER`
- You can view the `REFRESH_METHOD` in the `ALL_MVIEWS` data dictionary view.



- Manual refresh
  - Specify `ON DEMAND` option
  - By using the `DBMS_MVIEW` package
- Automatic refresh Synchronous
  - Specify `ON COMMIT` option
  - Upon commit of changes to the underlying tables but independent of the committing transaction
- Automatic refresh Asynchronous
  - Specify using `START WITH` and `NEXT` clauses
  - Defines a refresh interval for the materialized view
- `REFRESH_MODE` in `ALL_MVIEWS`



## Manual Refresh with `DBMS_MVIEW`

- For `ON DEMAND` refresh only
- Three procedures with the `DBMS_MVIEW` package:
  - `REFRESH`
  - `REFRESH_ALL_MVIEWS`
  - `REFRESH_DEPENDENT`



Specific materialized views:

```
Exec DBMS_MVIEW.REFRESH('cust_sales_mv');
```

### Materialized views based on one or more tables:

```
VARIABLE fail NUMBER;  
exec DBMS_MVIEW.REFRESH_DEPENDENT(-  
:fail, 'CUSTOMERS, SALES');
```

### All materialized views due for refresh:

```
VARIABLE fail NUMBER;  
exec DBMS_MVIEW.REFRESH_ALL_MVIEWS(:fail);
```



- If you want to use a materialized view instead of the base tables, a query must be rewritten.
- Query rewrites are transparent to applications.
- Query rewrites do not require special privileges on the materialized view.
- A materialized view can be enabled or disabled for query rewrites.



- Use `EXPLAIN PLAN` or `AUTOTRACE` to verify that query rewrites occur.
- Check the query response:
  - Fewer blocks are accessed.
  - Response time should be significantly better.

## Enabling and Controlling Query Rewrites



- Query rewrites are available with cost-based optimization only.

- The following parameters control query rewrites:

```
QUERY_REWRITE_ENABLED = {true|false|force}  
QUERY_REWRITE_INTEGRITY =  
{enforced|trusted|stale_tolerated}
```



```
EXPLAIN PLAN FOR
SELECT    t.week_ending_day
,         t.calendar_year
,         p.prod_subcategory
,         sum(c.unit_cost) AS dollars
FROM      costs c
,         times t
,         products p
WHERE     c.time_id = t.time_id
...

```

### Execution Plan

```
-----
0      SELECT STATEMENT Optimizer=ALL_ROWS (Cost...)
1      0      MAT_VIEW REWRITE ACCESS (FULL) OF 'costs_per_year_mv' (
              MAT_VIEW REWRITE) (Cost...)


```





```
SELECT    t.week_ending_day
,         t.calendar_year
,         p.prod_subcategory
,         sum(c.unit_cost) AS dollars
FROM      costs c, times t, products p
WHERE     c.time_id = t.time_id
AND       c.prod_id = p.prod_id
AND       t.calendar_year = '1999'
GROUP BY  t.week_ending_day, t.calendar_year
,         p.prod_subcategory
HAVING    sum(c.unit_cost) > 10000;
```

```
SELECT    week_ending_day
,         prod_subcategory
,         dollars
FROM      cost_per_year_mv
WHERE     calendar_year = '1999'
AND       dollars > 10000;
```

A blue arrow pointing downwards from the first query block to the second query block.



```
CREATE MATERIALIZED VIEW cust_orders_mv
ENABLE QUERY REWRITE AS
SELECT c.customer_id, SUM(order_total) AS amt
FROM   oe.orders s, oe.customers c
WHERE  s.customer_id = c.customer_id
GROUP BY c.customer_id;
```



```
SELECT /*+ REWRITE_OR_ERROR */ c.customer_id,
SUM(order_total) AS amt
FROM   oe.orders s, oe.customers c
WHERE  s.customer_id = c.customer_id
GROUP BY c.customer_id;
```



```
ORA-30393: a query block in the statement did
not rewrite
```



- For a given workload, the SQL Access Advisor:
  - Recommends creating the appropriate:
    - Materialized views
    - Materialized view logs
    - Indexes
  - Provides recommendations to optimize for:
    - Fast refresh
    - Query rewrite
  - Can be run:
    - From Oracle Enterprise Manager by using the SQL Access Advisor Wizard
    - By invoking the `DBMS_ADVISOR` package

# Using the DBMS\_MVIEW Package



## ➤ DBMS\_MVIEW methods

- EXPLAIN\_MVIEW
- EXPLAIN\_REWRITE
- TUNE\_MVIEW



## Tuning Materialized Views for Fast Refresh and Query Rewrite

```
DBMS_ADVISOR.TUNE_MVIEW (  
    task_name IN OUT VARCHAR2,  
    mv_create_stmt IN [CLOB | VARCHAR2]  
);
```

## Results of Tune\_MVIEW



- IMPLEMENTATION recommendations
  - CREATE MATERIALIZED VIEW LOG statements
  - ALTER MATERIALIZED VIEW LOG FORCE statements
  - One or more CREATE MATERIALIZED VIEW statements
- UNDO recommendations
  - DROP MATERIALIZED VIEW statements



## DBMS\_MVIEW.EXPLAIN\_MVIEW Procedure

- Accepts:
  - Materialized view name
  - SQL statement
- Advises what is and what is not possible:
  - For an existing materialized view
  - For a potential materialized view before you create it
- Stores results in `MV_CAPABILITIES_TABLE` (relational table) or in a `VARRAY`
- `utlxml.sql` must be executed as the current user to create `MV_CAPABILITIES_TABLE`.

```
EXEC dbms_mview.explain_mvview (
    'cust_sales_mv', '123');
```

```
SELECT capability_name, possible,  
related_text,msgtxt  
FROM mv_capabilities table
```

CAPABILITY	NAME	P	RELATED	TE	MSGT	TEXT
REFRESH	COMPLETE	Y				
REFRESH	FAST	N				
REWRITE		N				
PCT	TABLE	N	SALES		no	
partition	key					
or						
						PMARKER
in	select					

in select		
PCT TABLE	N CUSTOMERS	list relation





## ➤ Query rewrite considerations:

- Constraints
- Outer joins
- Text match
- Aggregates
- Grouping conditions
- Expression matching
- Date folding
- Statistics



<b>REWRITE</b>	<b>Rewrites a query in terms of materialized views</b>
<b>REWRITE_OR_ERROR</b>	<b>Forces an error if a query rewrite is not possible</b>
<b>NO_REWRITE</b>	<b>Disables query rewrite for the query block</b>

# SUMMARY

- In this lesson, you should have learned how to:
  - Create materialized views
  - Enable query rewrites using materialized views

# SUMMARY

- » Introduction to SQL tuning
- » Describe why the SQL statements are performing poorly
- » Introduction to Oracle Optimizer
- » Discuss the need for Optimizer
- » Explain the various phases of Optimization
- » Gather Execution Plans
- » Interpret Execution Plans
- » Interpret the output of TKPROF
- » Gather Optimizer statistics
- » Use Hints appropriately